

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [7 February - 13 February](#) / [Quiz-1: 10 AM](#)

**Started on** Saturday, 12 February 2022, 10:00:12 AM

**State** Finished

**Completed on** Saturday, 12 February 2022, 12:00:12 PM

**Time taken** 2 hours

**Grade** 7.58 out of 10.00 (76%)

Question **1**

Complete

Mark 0.30 out of 0.50

Select Yes if the mentioned element should be a part of PCB

Select No otherwise.

Yes	No	
<input checked="" type="radio"/>	<input type="radio"/>	Pointer to the parent process
<input checked="" type="radio"/>	<input type="radio"/>	PID of Init
<input checked="" type="radio"/>	<input type="radio"/>	List of opened files
<input type="radio"/>	<input checked="" type="radio"/>	Process context
<input type="radio"/>	<input checked="" type="radio"/>	EIP at the time of context switch
<input checked="" type="radio"/>	<input type="radio"/>	PID
<input checked="" type="radio"/>	<input type="radio"/>	Process state
<input checked="" type="radio"/>	<input type="radio"/>	Function pointers to all system calls
<input type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT
<input checked="" type="radio"/>	<input type="radio"/>	Memory management information about that process

Pointer to the parent process: Yes

PID of Init: No

List of opened files: Yes

Process context: Yes

EIP at the time of context switch: Yes

PID: Yes

Process state: Yes

Function pointers to all system calls: No

Pointer to IDT: No

Memory management information about that process: Yes

Question **2**

Complete

Mark 0.50 out of 0.50

Consider the following programs

`exec1.c`

```
#include <unistd.h>
#include <stdio.h>
int main() {
    execl("./exec2", "./exec2", NULL);
}
```

`exec2.c`

```
#include <unistd.h>
#include <stdio.h>
int main() {
    execl("/bin/ls", "/bin/ls", NULL);
    printf("hello\n");
}
```

Compiled as

```
cc exec1.c -o exec1
cc exec2.c -o exec2
```

And run as

`$ ./exec1`

Explain the output of the above command (`./exec1`)

Assume that `/bin/ls`, i.e. the 'ls' program exists.

Select one:

- ☐ a. Execution fails as the call to `execl()` in `exec1` fails
- ☐ b. Execution fails as one exec can't invoke another exec
- ☐ c. Execution fails as the call to `execl()` in `exec2` fails
- ☐ d. Program prints hello
- ☒ e. "ls" runs on current directory

The correct answer is: "ls" runs on current directory

Question **3**

Complete

Mark 0.50 out of  
0.50

What's the trapframe in xv6?

- ☐ a. A frame of memory that contains all the trap handler code's function pointers
- ☐ b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only
- ☐ c. The IDT table
- ☒ d. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S
- ☐ e. A frame of memory that contains all the trap handler code
- ☐ f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- ☐ g. A frame of memory that contains all the trap handler's addresses

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question **4**

Complete

Mark 0.33 out of 0.50

The `bootmain()` function has this code

```
elf = (struct elfhdr*)0x10000; // scratch space
readseg((uchar*)elf, 4096, 0);
```

Mark the statements as True or False with respect to this code.

In these statements 0x1000 is referred to as ADDRESS

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	It the value of ADDRESS is changed to a higher number (upto a limit), the program could still work
<input type="radio"/>	<input checked="" type="radio"/>	The value ADDRESS is changed to a 0 the program could still work
<input type="radio"/>	<input checked="" type="radio"/>	If the value of ADDRESS is changed, then the program will not work
<input checked="" type="radio"/>	<input type="radio"/>	This line effectively loads the ELF header and the program headers at ADDRESS
<input checked="" type="radio"/>	<input type="radio"/>	This line loads the kernel code at ADDRESS
<input checked="" type="radio"/>	<input type="radio"/>	It the value of ADDRESS is changed to a lower number (upto a limit), the program could still work

It the value of ADDRESS is changed to a higher number (upto a limit), the program could still work: True

The value ADDRESS is changed to a 0 the program could still work: False

If the value of ADDRESS is changed, then the program will not work: False

This line effectively loads the ELF header and the program headers at ADDRESS: False

This line loads the kernel code at ADDRESS: False

It the value of ADDRESS is changed to a lower number (upto a limit), the program could still work: True

Question **5**

Complete

Mark 0.50 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

Process P1 is running

Process P0 is running

context of P1 is loaded from P1's PCB

Control is passed to P1

context of P0 is saved in P0's PCB

timer interrupt occurs

The correct answer is: Process P1 is running → 6, Process P0 is running → 1, context of P1 is loaded from P1's PCB → 4, Control is passed to P1 → 5, context of P0 is saved in P0's PCB → 3, timer interrupt occurs → 2

Question **6**

Complete

Mark 0.07 out of 0.50

Order the events that occur on a timer interrupt:

Execute the code of the new process

Save the context of the currently running process

Select another process for execution

Jump to a code pointed by IDT

Set the context of the new process

Change to kernel stack of currently running process

Jump to scheduler code

The correct answer is: Execute the code of the new process → 7, Save the context of the currently running process → 3, Select another process for execution → 5, Jump to a code pointed by IDT → 2, Set the context of the new process → 6, Change to kernel stack of currently running process → 1, Jump to scheduler code → 4

Question **7**

Complete

Mark 0.50 out of 0.50

In bootasm.S, on the line

```
ljmp    $(SEG_KCODE<<3), $start32
```

The SEG\_KCODE << 3, that is shifting of 1 by 3 bits is done because

- ☐ a. The value 8 is stored in code segment
- ☒ b. The code segment is 16 bit and only upper 13 bits are used for segment number
- ☐ c. The ljmp instruction does a divide by 8 on the first argument
- ☐ d. The code segment is 16 bit and only lower 13 bits are used for segment number
- ☐ e. While indexing the GDT using CS, the value in CS is always divided by 8

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question **8**

Complete

Mark 0.50 out of 0.50

Suppose a program does a scanf() call.

Essentially the scanf does a read() system call.

This call will obviously "block" waiting for the user input.

In terms of OS data structures and execution of code, what does it mean?

Select one:

- ☐ a. read() will return and process will be taken to a wait queue
- ☐ b. OS code for read() will move the PCB of this process to a wait queue and return from the system call
- ☒ c. OS code for read() will move PCB of current process to a wait queue and call scheduler
- ☐ d. read() returns and process calls scheduler()
- ☐ e. OS code for read() will call scheduler

The correct answer is: OS code for read() will move PCB of current process to a wait queue and call scheduler

## Question 9

Complete

Mark 0.67 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

(Note: non-interruptible kernel code means, if the kernel code is executing, then interrupts will be disabled).

Note: A possible sequence may have some missing steps in between. An impossible sequence will have n and n+1th steps such that n+1th step can not follow n'th step.

Select one or more:

- ☐ a. P1 running  
keyboard hardware interrupt  
keyboard interrupt handler running  
interrupt handler returns  
P1 running  
P1 makes system call  
system call returns  
P1 running  
timer interrupt  
scheduler  
P2 running
- ☐ b. P1 running  
P1 makes system call and blocks  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P3 running  
Hardware interrupt  
Interrupt unblocks P1  
Interrupt returns  
P3 running  
Timer interrupt  
Scheduler  
P1 running
- ☒ c.  
P1 running  
P1 makes system call  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P1 running again
- ☒ d. P1 running  
P1 makes system call  
timer interrupt  
Scheduler  
P2 running  
timer interrupt  
Scheduler  
P1 running  
P1's system call return
- ☐ e. P1 running  
P1 makes system call  
system call returns  
P1 running

timer interrupt  
Scheduler running  
P2 running

- ☐ f. P1 running  
P1 makes system call and blocks  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P1 running again

The correct answers are: P1 running  
P1 makes system call and blocks  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P1 running again, P1 running  
P1 makes system call  
timer interrupt  
Scheduler  
P2 running  
timer interrupt  
Scheduler  
P1 running  
P1's system call return,  
P1 running  
P1 makes system call  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P1 running again



Question **10**

Complete

Mark 0.75 out of  
1.00

Select the correct statements about interrupt handling in xv6 code

- ☒ a. All the 256 entries in the IDT are filled
- ☒ b. The function `trap()` is the called irrespective of hardware interrupt/system-call/exception
- ☐ c. xv6 uses the 0x64th entry in IDT for system calls
- ☐ d. On any interrupt/syscall/exception the control first jumps in `trapasm.S`
- ☒ e. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- ☒ f. xv6 uses the 64th entry in IDT for system calls
- ☐ g. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- ☐ h. The CS and EIP are changed only immediately on a hardware interrupt
- ☐ i. Before going to `alltraps`, the kernel stack contains upto 5 entries.
- ☐ j. The `trapframe` pointer in `struct proc`, points to a location on user stack
- ☒ k. On any interrupt/syscall/exception the control first jumps in `vectors.S`
- ☒ l. The `trapframe` pointer in `struct proc`, points to a location on kernel stack
- ☐ m. The function `trap()` is the called only in case of hardware interrupt

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in `vectors.S`, Before going to `alltraps`, the kernel stack contains upto 5 entries., The `trapframe` pointer in `struct proc`, points to a location on kernel stack, The function `trap()` is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question **11**

Complete

Mark 0.50 out of 0.50

Select all the correct statements about zombie processes

Select one or more:

- ☒ a. A process can become zombie if it finishes, but the parent has finished before it
- ☐ b. A process becomes zombie when it's parent finishes
- ☒ c. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it
- ☐ d. A zombie process remains zombie forever, as there is no way to clean it up
- ☐ e. Zombie processes are harmless even if OS is up for long time
- ☒ f. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent
- ☒ g. A zombie process occupies space in OS data structures
- ☒ h. init() typically keeps calling wait() for zombie processes to get cleaned up

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

Question **12**

Complete

Mark 0.50 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- ☒ a. The setting up of the most essential memory management infrastructure needs assembly code
- ☐ b. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time
- ☒ c. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C
- ☐ d. The code for reading ELF file can not be written in assembly

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question **13**

Complete

Mark 0.50 out of 0.50

For each line of code mentioned on the left side, select the location of sp/esp that is in use

`readseg((uchar*)elf, 4096, 0);`  
in `bootmain.c`

0x7c00 to 0

`jmp *%eax`  
in `entry.S`

The 4KB area in kernel image, loaded in memory, named as 'stack'

`ljmp $(SEG_KCODE<<3), $start32`  
in `bootasm.S`

Immaterial as the stack is not used here

`cli`

in `bootasm.S`

Immaterial as the stack is not used here

`call bootmain`  
in `bootasm.S`

0x7c00 to 0

The correct answer is: `readseg((uchar*)elf, 4096, 0);`

in `bootmain.c` → 0x7c00 to 0, `jmp *%eax`

in `entry.S` → The 4KB area in kernel image, loaded in memory, named as 'stack', `ljmp $(SEG_KCODE<<3), $start32`

in `bootasm.S` → Immaterial as the stack is not used here, `cli`

in `bootasm.S` → Immaterial as the stack is not used here, `call bootmain`

in `bootasm.S` → 0x7c00 to 0

## Question 14

Complete

Mark 0.80 out of 1.00

Mark the statements, w.r.t. the scheduler of xv6 as True or False

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The function <code>scheduler()</code> executes using the kernel-only stack
<input type="radio"/>	<input checked="" type="radio"/>	The work of selecting and scheduling a process is done only in <code>scheduler()</code> and not in <code>sched()</code>
<input type="radio"/>	<input checked="" type="radio"/>	the control returns to <code>switchkvm();</code> after <code>swtch(&amp;(c-&gt;scheduler), p-&gt;context);</code> in <code>scheduler()</code>
<input checked="" type="radio"/>	<input type="radio"/>	<code>swtch</code> is a function that saves old context, loads new context, and returns to last EIP in the new context
<input checked="" type="radio"/>	<input type="radio"/>	<code>swtch</code> is a function that does not return to the caller
<input checked="" type="radio"/>	<input type="radio"/>	The variable <code>c-&gt;scheduler</code> on first processor uses the stack allocated entry.S
<input checked="" type="radio"/>	<input type="radio"/>	<code>sched()</code> and <code>scheduler()</code> are co-routines
<input checked="" type="radio"/>	<input type="radio"/>	When a process is scheduled for execution, it resumes execution in <code>sched()</code> after the call to <code>swtch()</code>
<input type="radio"/>	<input checked="" type="radio"/>	the control returns to <code>mycpu()-&gt;intena = intena; ();</code> after <code>swtch(&amp;p-&gt;context, mycpu()-&gt;scheduler);</code> in <code>sched()</code>
<input checked="" type="radio"/>	<input type="radio"/>	<code>sched()</code> calls <code>scheduler()</code> and <code>scheduler()</code> calls <code>sched()</code>

The function `scheduler()` executes using the kernel-only stack: TrueThe work of selecting and scheduling a process is done only in `scheduler()` and not in `sched()`: Truethe control returns to `switchkvm();` after `swtch(&(c->scheduler), p->context);` in `scheduler()`: False`swtch` is a function that saves old context, loads new context, and returns to last EIP in the new context: True`swtch` is a function that does not return to the caller: TrueThe variable `c->scheduler` on first processor uses the stack allocated entry.S: True`sched()` and `scheduler()` are co-routines: TrueWhen a process is scheduled for execution, it resumes execution in `sched()` after the call to `swtch()`: True

the control returns to `mycpu()->intena = intena; ();` after `swtch(&p->context, mycpu()->scheduler);` in `sched()`: **False**  
`sched()` **calls** `scheduler()` and `scheduler()` **calls** `sched()`: **False**

Question **15**

Complete

Mark 0.50 out of 1.00

Which parts of the xv6 code in `bootasm.S`, `bootmain.c`, `entry.S` and in the codepath related to `scheduler()` and `trap handling()` can also be written in some other way, and still ensure that xv6 works properly?

Writing code is not necessary. You only need to comment on which part of the code could be changed to something else or written in another fashion.

Maximum two points to be written.

1)we can change the scheduling algorithm

2)Still xv6 will work

Question **16**

Complete

Mark 0.16 out of 0.50

Select all the correct statements about code of bootmain() in xv6

```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}
```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- ☒ a. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded
- ☐ b. The stosb() is used here, to fill in some space in memory with zeroes
- ☒ c. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it.
- ☐ d. The condition if(ph->memsz > ph->filesz) is never true.
- ☒ e. The elf->entry is set by the linker in the kernel file and it's 8010000c
- ☐ f. The readseg finally invokes the disk I/O code using assembly instructions
- ☒ g. The kernel file gets loaded at the Physical address 0x10000 in memory.
- ☐ h. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- ☒ i. The elf->entry is set by the linker in the kernel file and it's 0x80000000

- ☐ j. The kernel file has only two program headers
- ☐ k. The kernel file gets loaded at the Physical address 0x10000 + 0x80000000 in memory.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

◀ [Extra Reading on Linkers: A writeup by Ian Taylor \(keep changing url string from 38 to 39, and so on\)](#)

Jump to...

[\(Code\) IPC - Shm, Messages](#) ▶

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-20-21](#) / [OS-Even-sem-2020-21](#) / 14 February - 20 February  
/ [Quiz-1](#)

**Started on** Saturday, 20 February 2021, 2:51 PM

**State** Finished

**Completed on** Saturday, 20 February 2021, 3:55 PM

**Time taken** 1 hour 3 mins

**Grade** 7.30 out of 20.00 (37%)

#### Question 1

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements about the state of a process.

- ☒ a. A process can self-terminate only when it's running ✓
- ☒ b. Typically, it's represented as a number in the PCB ✓
- ☒ c. A process that is running is not on the ready queue ✓
- ☒ d. Processes in the ready queue are in the ready state ✓
- ☐ e. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers
- ☒ f. Changing from running state to waiting state results in "giving up the CPU" ✓
- ☐ g. A process in ready state is ready to receive interrupts
- ☒ h. A waiting process starts running after the wait is over ✗
- ☒ i. A process changes from running to ready state on a timer interrupt ✓
- ☒ j. A process in ready state is ready to be scheduled ✓
- ☒ k. A running process may terminate, or go to wait or become ready again ✓
- ☒ l. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓
- ☐ m. A process waiting for any condition is woken up by another process only
- ☐ n. A process changes from running to ready state on a timer interrupt or any I/O wait

Your answer is partially correct.

You have selected too many options.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code



## Question 2

Incorrect

Mark 0.00 out of 1.00

For each line of code mentioned on the left side, select the location of sp/esp that is in use

<code>jmp *%eax</code> in <code>entry.S</code>	0x7c00 to 0x10000	✗
<code>ljmp \$(SEG_KCODE&lt;&lt;3), \$start32</code> in <code>bootasm.S</code>	0x10000 to 0x7c00	✗
<code>call bootmain</code> in <code>bootasm.S</code>	0x7c00 to 0x10000	✗
<code>cli</code> in <code>bootasm.S</code>	0x7c00 to 0	✗
<code>readseg((uchar*)elf, 4096, 0);</code> in <code>bootmain.c</code>	The 4KB area in kernel image, loaded in memory, named as 'stack'	✗

Your answer is incorrect.

The correct answer is: `jmp *%eax`

in `entry.S` → The 4KB area in kernel image, loaded in memory, named as 'stack', `ljmp $(SEG_KCODE<<3), $start32`

in `bootasm.S` → Immaterial as the stack is not used here, `call bootmain`

in `bootasm.S` → 0x7c00 to 0, `cli`

in `bootasm.S` → Immaterial as the stack is not used here, `readseg((uchar*)elf, 4096, 0);`

in `bootmain.c` → 0x7c00 to 0

## Question 3

Correct

Mark 0.25 out of 0.25

Order the following events in boot process (from 1 onwards)

Boot loader	2	✓
Shell	6	✓
BIOS	1	✓
OS	3	✓
Init	4	✓
Login interface	5	✓

Your answer is correct.

The correct answer is: Boot loader → 2, Shell → 6, BIOS → 1, OS → 3, Init → 4, Login interface → 5

## Question 4

Partially correct

Mark 0.30 out of 0.50

Consider the following command and it's output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
    dd if=/dev/zero of=xv6.img count=10000
    dd if=bootblock of=xv6.img conv=notrunc
    dd if=kernel of=xv6.img seek=1 conv=notrunc
```

```
kernel: $(OBJJS) entry.o entryother initcode kernel.ld
    $(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJJS) -b binary initcode entryother
    $(OBJDUMP) -S kernel > kernel.asm
    $(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- ☐ a. Although the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- ☒ b. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- ☒ c. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- ☐ d. The bootmain() code does not read the kernel completely in memory
- ☐ e. readseg() reads first 4k bytes of kernel in memory
- ☐ f. Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.
- ☐ g. The kernel.asm file is the final kernel file
- ☐ h. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not read as it is user programs.
- ☒ i. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain()., readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question **5**

Partially correct

Mark 0.50 out of 1.00

```
int f() {  
    int count;  
    for (count = 0; count < 2; count++) {  
        if (fork() == 0)  
            printf("Operating-System\n");  
    }  
    printf("TYCOMP\n");  
}
```

The number of times "Operating-System" is printed, is:

Answer:  ☒

The correct answer is: 7.00

## Question 6

Partially correct

Mark 0.40 out of 0.50

Select Yes/True if the mentioned element must be a part of PCB

Select No/False otherwise.

Yes	No		
<input checked="" type="radio"/>	<input type="radio"/>	PID	✓
<input checked="" type="radio"/>	<input type="radio"/>	Process context	✓
<input checked="" type="radio"/>	<input type="radio"/>	List of opened files	✓
<input checked="" type="radio"/>	<input type="radio"/>	Process state	✓
<input type="radio"/>	<input checked="" type="radio"/>	Parent's PID	✗
<input type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT	✓
<input type="radio"/>	<input checked="" type="radio"/>	Function pointers to all system calls	✓
<input checked="" type="radio"/>	<input type="radio"/>	Memory management information about that process	✓
<input checked="" type="radio"/>	<input type="radio"/>	Pointer to the parent process	✗
<input checked="" type="radio"/>	<input type="radio"/>	EIP at the time of context switch	✓

PID: Yes

Process context: Yes

List of opened files: Yes

Process state: Yes

Parent's PID: No

Pointer to IDT: No

Function pointers to all system calls: No

Memory management information about that process: Yes

Pointer to the parent process: Yes

EIP at the time of context switch: Yes

## Question 7

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about code of bootmain() in xv6

```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*) (void))(elf->entry);
    entry();
}
```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- ☒ a. The kernel file gets loaded at the Physical address 0x10000 + 0x80000000 in memory. ✗
- ☒ b. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- ☒ c. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded ✓
- ☒ d. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it. ✓
- ☒ e. The kernel file has only two program headers ✓
- ☒ f. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- ☒ g. The readseg finally invokes the disk I/O code using assembly instructions ✓
- ☒ h. The elf->entry is set by the linker in the kernel file and it's 8010000c ✓
- ☒ i. The kernel file gets loaded at the Physical address 0x10000 in memory. ✓
- ☒ j. The condition if(ph->memsz > ph->filesz) is never true. ✗
- ☒ k. The stosb() is used here, to fill in some space in memory with zeroes ✓

Your answer is incorrect.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

**Question 8**

Partially correct

Mark 0.13 out of 0.25

Which of the following are NOT a part of job of a typical compiler?

- ☒ a. Check the program for logical errors
- ☐ b. Convert high level language code to machine code
- ☐ c. Process the # directives in a C program
- ☐ d. Invoke the linker to link the function calls with their code, extern globals with their declaration
- ☐ e. Check the program for syntactical errors
- ☐ f. Suggest alternative pieces of code that can be written



Your answer is partially correct.

You have correctly selected 1.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

**Question 9**

Correct

Mark 0.25 out of 0.25

Rank the following storage systems from slowest (first) to fastest(last)

Cache	6	✓
Hard Disk	3	✓
RAM	5	✓
Optical Disks	2	✓
Non volatile memory	4	✓
Registers	7	✓
Magnetic Tapes	1	✓

Your answer is correct.

The correct answer is: Cache → 6, Hard Disk → 3, RAM → 5, Optical Disks → 2, Non volatile memory → 4, Registers → 7, Magnetic Tapes → 1

Question **10**

Partially correct

Mark 0.21 out of 0.50

Which of the following parts of a C program do not have any corresponding machine code ?

- ☐ a. local variable declaration
- ☐ b. global variables
- ☒ c. function calls
- ☒ d. #directives
- ☐ e. expressions
- ☐ f. pointer dereference
- ☒ g. typedefs

✗

✓

✓

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: #directives, typedefs, global variables

Question **11**

Correct

Mark 0.25 out of 0.25

Match a system call with it's description

pipe	create an unnamed FIFO storage with 2 ends - one for reading and another for writing	✓
dup	create a copy of the specified file descriptor into smallest available file descriptor	✓
dup2	create a copy of the specified file descriptor into another specified file descriptor	✓
exec	execute a binary file overlaying the image of current process	✓
fork	create an identical child process	✓

Your answer is correct.

The correct answer is: pipe → create an unnamed FIFO storage with 2 ends - one for reading and another for writing, dup → create a copy of the specified file descriptor into smallest available file descriptor, dup2 → create a copy of the specified file descriptor into another specified file descriptor, exec → execute a binary file overlaying the image of current process, fork → create an identical child process

Question **12**

Correct

Mark 0.25 out of 0.25

Match the register with the segment used with it.

eip	cs	✓
edi	es	✓
esi	ds	✓
ebp	ss	✓
esp	ss	✓

Your answer is correct.

The correct answer is: eip → cs, edi → es, esi → ds, ebp → ss, esp → ss

Question **13**

Correct

Mark 0.25 out of 0.25

What's the trapframe in xv6?

- ☐ a. A frame of memory that contains all the trap handler code
- ☐ b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- ☐ c. The IDT table
- ☐ d. A frame of memory that contains all the trap handler code's function pointers
- ☐ e. A frame of memory that contains all the trap handler's addresses
- ☒ f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S ✓
- ☐ g. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only

Your answer is correct.

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S



## Question 14

Incorrect

Mark 0.00 out of 0.50

Select all the correct statements about linking and loading.

Select one or more:

- ☒ a. Continuous memory management schemes can support dynamic linking and dynamic loading. ✗
- ☒ b. Loader is last stage of the linker program ✗
- ☒ c. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently) ✓
- ☒ d. Dynamic linking and loading is not possible without demand paging or demand segmentation. ✓
- ☒ e. Dynamic linking essentially results in relocatable code. ✓
- ☒ f. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- ☒ g. Loader is part of the operating system ✓
- ☒ h. Static linking leads to non-relocatable code ✗
- ☒ i. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

## Question 15

Incorrect

Mark 0.00 out of 0.25

In bootasm.S, on the line

```
ljmp $(SEG_KCODE<<3), $start32
```

The SEG\_KCODE << 3, that is shifting of 1 by 3 bits is done because

- ☐ a. The value 8 is stored in code segment
- ☐ b. The code segment is 16 bit and only upper 13 bits are used for segment number
- ☒ c. The code segment is 16 bit and only lower 13 bits are used for segment number ✗
- ☐ d. While indexing the GDT using CS, the value in CS is always divided by 8
- ☐ e. The ljmp instruction does a divide by 8 on the first argument

Your answer is incorrect.

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question **16**

Partially correct

Mark 0.07 out of 0.50

Order the events that occur on a timer interrupt:

Change to kernel stack	1	✗
Jump to a code pointed by IDT	2	✗
Jump to scheduler code	5	✗
Set the context of the new process	4	✗
Save the context of the currently running process	3	✓
Execute the code of the new process	6	✗
Select another process for execution	7	✗

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: Change to kernel stack → 2, Jump to a code pointed by IDT → 1, Jump to scheduler code → 4, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7, Select another process for execution → 5

Question **17**

Incorrect

Mark 0.00 out of 1.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

**\$ ls ./tmp/asdfksdf >/tmp/ddd 2>&1**

**Program 1**

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    dup(fd);
    close(2);
    dup(fd);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

**Program 2**

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    close(1);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(2);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Select all the correct statements about the programs

Select one or more:

- ☒ a. Both programs are correct ✗
- ☒ b. Program 2 makes sure that there is one file offset used for '2' and '1' ✗
- ☒ c. Only Program 2 is correct ✗
- ☒ d. Program 2 does 1>&2 ✗
- ☒ e. Program 2 ensures 2>&1 and does not ensure > /tmp/ddd ✗
- ☒ f. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- ☒ g. Program 1 is correct for > /tmp/ddd but not for 2>&1 ✗
- ☒ h. Program 1 does 1>&2 ✗
- ☒ i. Both program 1 and 2 are incorrect ✗
- ☒ j. Program 2 is correct for > /tmp/ddd but not for 2>&1 ✗
- ☒ k. Only Program 1 is correct ✓
- ☒ l. Program 1 ensures 2>&1 and does not ensure > /tmp/ddd ✗

Your answer is incorrect.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'



Question **18**

Correct

Mark 0.25 out of 0.25

Select the option which best describes what the CPU does during it's powered ON lifetime

- ☐ a. Ask the user what is to be done, and execute that task
- ☐ b. Ask the OS what is to be done, and execute that task
- ☐ c. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask the User or the OS what is to be done next, repeat
- ☒ d. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat ✓
- ☐ e. Fetch instruction specified by OS, Decode and execute it, repeat
- ☐ f. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask OS what is to be done next, repeat

The correct answer is: Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat

## Question 19

Partially correct

Mark 0.86 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code.

```
int main(int argc, char *argv[]) {  
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;  
  
    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    fd2 = open("/tmp/2", O_RDONLY);  
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    close(0);  
    close(1);  
    dup(fd2);  
    dup(fd3);  
    close(fd3);  
    dup2(fd2, fd4);  
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);  
    return 0;  
}
```

1	<input type="text" value="closed"/>	✖
fd4	<input type="text" value="/tmp/2"/>	✔
fd2	<input type="text" value="/tmp/2"/>	✔
fd1	<input type="text" value="/tmp/1"/>	✔
2	<input type="text" value="stderr"/>	✔
0	<input type="text" value="/tmp/2"/>	✔
fd3	<input type="text" value="closed"/>	✔

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: 1 → /tmp/3, fd4 → /tmp/2, fd2 → /tmp/2, fd1 → /tmp/1, 2 → stderr, 0 → /tmp/2, fd3 → closed

Question **20**

Incorrect

Mark 0.00 out of 2.00

Following code claims to implement the command

```
/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1
```

Fill in the blanks to make the code work.

**Note:** Do not include space in writing any option. `x[1][2]` should be written without any space, and so is the case with `[1]` or `[2]`. Pay attention to exact syntax and do not write any extra character like `'` or `=` etc.

```
int main(int argc, char *argv[]) {
```

```
    int pid1, pid2;
```

```
    int pfd[
```

```
    x ][2];
```

```
    pipe(
```

```
    x );
```

```
    pid1 =
```

```
    x ;
```

```
    if(pid1 != 0) {
```

```
        close(pfd[0]
```

```
    x );
```

```
        close(
```

```
    x );
```

```
        dup(
```

```
    x );
```

```
        execl("/bin/ls", "/bin/ls", "
```

```
    x ", NULL);
```

```
    }
```

```
    pipe(
```

```
    x );
```

```
    x = fork();
```

```
    if(pid2 == 0) {
```

```
        close(
```

```
    x ;
```

```
        close(0);
```

```
        dup(
```

```
    x );
```

```
        close(pfd[1]
```

```
    × );  
        close(  
              
        × );  
        dup(  
              
        × );  
        execl("/usr/bin/head", "/usr/bin/head", "  
              
        × ", NULL);  
    } else {  
        close(pfd  
              
        × );  
        close(  
              
        × );  
        dup(  
              
        × );  
        close(pfd  
              
        × );  
        execl("/usr/bin/tail", "/usr/bin/tail", "  
              
        × ", NULL);  
    }  
}
```

Question **21**

Partially correct

Mark 0.11 out of 1.00

Select all the correct statements about calling convention on x86 32-bit.

- ☒ a. Return address is one location above the ebp ✓
- ☒ b. Parameters may be passed in registers or on stack ✓
- ☒ c. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function ✓
- ☒ d. The ebp pointers saved on the stack constitute a chain of activation records ✓
- ☒ e. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables ✗
- ☒ f. Parameters may be passed in registers or on stack ✓
- ☒ g. The return value is either stored on the stack or returned in the eax register ✗
- ☐ h. Paramters are pushed on the stack in left-right order
- ☐ i. during execution of a function, ebp is pointing to the old ebp
- ☒ j. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function ✗
- ☒ k. Compiler may allocate more memory on stack than needed ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question **22**

Correct

Mark 1.00 out of 1.00

Match the program with it's output (ignore newlines in the output. Just focus on the count of the number of 'hi')

- |   |       |   |
|---|-------|---|
| <code>main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }</code> | hi    | ✓ |
| <code>main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }</code>                    | hi hi | ✓ |
| <code>main() { int i = NULL; fork(); printf("hi\n"); }</code>   | hi hi | ✓ |
| <code>main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }</code>                            | hi    | ✓ |

Your answer is correct.

The correct answer is: `main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi hi, `main() { int i = NULL; fork(); printf("hi\n"); }` → hi hi, `main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi



Question **23**

Incorrect

Mark 0.00 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?  
Select all the appropriate choices

- ☒ a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time ✗
- ☒ b. The setting up of the most essential memory management infrastructure needs assembly code ✓
- ☒ c. The code for reading ELF file can not be written in assembly ✗
- ☒ d. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C ✓

Your answer is incorrect.

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question **24**

Incorrect

Mark 0.00 out of 0.50

xv6.img: bootblock kernel

```
dd if=/dev/zero of=xv6.img count=10000
```

```
dd if=bootblock of=xv6.img conv=notrunc
```

```
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- ☒ a. The size of the kernel file is nearly 5 MB ✓
- ☒ b. The kernel is located at block-1 of the xv6.img ✗
- ☒ c. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk. ✗
- ☐ d. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)
- ☒ e. The bootblock is located on block-0 of the xv6.img ✗
- ☒ f. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk. ✓
- ☒ g. The bootblock may be 512 bytes or less (looking at the Makefile instruction) ✗
- ☒ h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file. ✗
- ☒ i. The size of the xv6.img is nearly 5 MB ✗
- ☒ j. xv6.img is the virtual processor used by the qemu emulator ✓
- ☒ k. Blocks in xv6.img after kernel may be all zeroes. ✗

Your answer is incorrect.

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question **25**

Incorrect

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

Select one or more:

- ☐ a. P1 running  
P1 makes system call  
timer interrupt  
Scheduler  
P2 running  
timer interrupt  
Scheduler  
P1 running  
P1's system call return
- ☒ b. P1 running  
P1 makes system call and blocks  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P1 running again
- ☐ c. P1 running  
P1 makes system call  
system call returns  
P1 running  
timer interrupt  
Scheduler running  
P2 running
- ☒ d. P1 running  
P1 makes system call and blocks  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P3 running  
Hardware interrupt  
Interrupt unblocks P1  
Interrupt returns  
P3 running  
Timer interrupt  
Scheduler  
P1 running
- ☐ e.  
P1 running  
P1 makes system call  
Scheduler  
P2 running  
P2 makes system call and blocks  
Scheduler  
P1 running again
- ☒ f. P1 running  
keyboard hardware interrupt  
keyboard interrupt handler running  
interrupt handler returns  
P1 running  
P1 makes system call  
system call returns



P1 running  
timer interrupt  
scheduler  
P2 running

Your answer is incorrect.

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheuler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question **26**

Correct

Mark 0.25 out of 0.25

Which of the following are the files related to bootloader in xv6?

- ☐ a. bootasm.s and entry.S
- ☒ b. bootasm.S and bootmain.c
- ☐ c. bootasm.S, bootmain.c and bootblock.c
- ☐ d. bootmain.c and bootblock.S



Your answer is correct.

The correct answer is: bootasm.S and bootmain.c

Question **27**

Correct

Mark 0.25 out of 0.25

Match the following parts of a C program to the layout of the process in memory

Instructions	Text section	✓
Local Variables	Stack Section	✓
Dynamically allocated memory	Heap Section	✓
Global and static data	Data section	✓

Your answer is correct.

The correct answer is:

Instructions → Text section, Local Variables → Stack Section,  
Dynamically allocated memory → Heap Section,  
Global and static data → Data section

Question **28**

Incorrect

Mark 0.00 out of 0.50

What will this program do?

```
int main() {  
    fork();  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("hello");  
}
```

- ☐ a. one process will run ls, another will print hello
- ☒ b. run ls once
- ☐ c. run ls twice
- ☐ d. run ls twice and print hello twice
- ☐ e. run ls twice and print hello twice, but output will appear in some random order

✗

Your answer is incorrect.

The correct answer is: run ls twice

Question **29**

Correct

Mark 0.25 out of 0.25

What is the OS Kernel?

- ☒ a. The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run ✔ correct
- ☐ b. The set of tools like compiler, linker, loader, terminal, shell, etc.
- ☐ c. Only the system programs like compiler, linker, loader, etc.
- ☐ d. Everything that I see on my screen

The correct answer is: The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run

Question **30**

Correct

Mark 0.50 out of 0.50

Which of the following is/are not saved during context switch?

- ☐ a. Program Counter
- ☐ b. General Purpose Registers
- ☒ c. Bus ✔
- ☐ d. Stack Pointer
- ☐ e. MMU related registers/information
- ☒ f. Cache ✔
- ☒ g. TLB ✔

Your answer is correct.

The correct answers are: TLB, Cache, Bus

## Question 31

Partially correct

Mark 0.10 out of 0.25

Select the order in which the various stages of a compiler execute.

Linking	3	✗
Syntactical Analysis	2	✓
Pre-processing	1	✓
Intermediate code generation	does not exist	✗
Loading	4	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Linking → 4, Syntactical Analysis → 2, Pre-processing → 1, Intermediate code generation → 3, Loading → does not exist

## Question 32

Partially correct

Mark 0.08 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

context of P0 is saved in P0's PCB	2	✗
context of P1 is loaded from P1's PCB	3	✗
Process P1 is running	5	✗
timer interrupt occurs	6	✗
Process P0 is running	1	✓
Control is passed to P1	4	✗

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: context of P0 is saved in P0's PCB → 3, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, timer interrupt occurs → 2, Process P0 is running → 1, Control is passed to P1 → 5

Question **33**

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- ☐ a. On any interrupt/syscall/exception the control first jumps in vectors.S
- ☐ b. The trapframe pointer in struct proc, points to a location on user stack
- ☐ c. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- ☐ d. xv6 uses the 64th entry in IDT for system calls
- ☐ e. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- ☐ f. The trapframe pointer in struct proc, points to a location on kernel stack
- ☐ g. The function trap() is the called only in case of hardware interrupt
- ☐ h. The CS and EIP are changed only immediately on a hardware interrupt
- ☐ i. All the 256 entries in the IDT are filled
- ☐ j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- ☐ k. The function trap() is the called irrespective of hardware interrupt/system-call/exception
- ☐ l. xv6 uses the 0x64th entry in IDT for system calls
- ☐ m. Before going to alltraps, the kernel stack contains upto 5 entries.

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

◀ (Assignment) [Change free list management in xv6](#)

Jump to...



[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [24 January - 30 January](#)  
/ [Topic-wise Quiz: 1](#)

**Started on** Monday, 24 January 2022, 8:42:38 PM

**State** Finished

**Completed on** Monday, 24 January 2022, 9:52:11 PM

**Time taken** 1 hour 9 mins

**Grade** 9.63 out of 20.00 (48%)

Question **1**

Complete

Mark 0.00 out of  
2.00

Which of the following instructions should be privileged?

Select one or more:

- ☒ a. Access memory management unit of the processor
- ☒ b. Turn off interrupts.
- ☒ c. Set value of a memory location
- ☒ d. Switch from user to kernel mode.
- ☒ e. Read the clock.
- ☒ f. Modify entries in device-status table
- ☒ g. Access a general purpose register
- ☒ h. Access I/O device.
- ☒ i. Set value of timer.

The correct answers are: Set value of timer., Access memory management unit of the processor, Turn off interrupts., Modify entries in device-status table, Access I/O device., Switch from user to kernel mode.

Question **2**

Complete

Mark 1.00 out of 1.00

Rank the following storage systems from slowest (first) to fastest(last)

You can drag and drop the items below/above each other.

Magnetic tapes

Optical disk

Hard-disk drives

Nonvolatile memory

Main memory

Cache

Registers

Question **3**

Complete

Mark 1.00 out of 1.00

```
int value = 5;
int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("%d", value); /* LINE A */
    }
    return 0;
}
```

What's the value printed here at LINE A?

Answer: 5

The correct answer is: 5

Question **4**

Complete

Mark 0.00 out of 0.50

Is the terminal a part of the kernel on GNU/Linux systems?

- ☐ a. no
- ☒ b. yes

The correct answer is: no

Question **5**

Complete

Mark 1.00 out of 1.00

Why should a program exist in memory before it starts executing ?

- ☐ a. Because the variables of the program are stored in memory
- ☐ b. Because the memory is volatile
- ☐ c. Because the hard disk is a slow medium
- ☒ d. Because the processor can run instructions and access data only from memory

The correct answer is: Because the processor can run instructions and access data only from memory

Question **6**

Complete

Mark 1.00 out of 1.00

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- ☒ a. It prohibits a user mode process from running privileged instructions
- ☐ b. It prohibits invocation of kernel code completely, if a user program is running
- ☐ c. It disallows hardware interrupts when a process is running
- ☐ d. It prohibits one process from accessing other process's memory

The correct answer is: It prohibits a user mode process from running privileged instructions

## Question 7

Complete

Mark 0.00 out of 2.00

Select all the correct statements about calling convention on x86 32-bit.

- ☒ a. Return address is one location above the ebp
- ☒ b. Parameters are pushed on the stack in left-right order
- ☒ c. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables
- ☒ d. Compiler may allocate more memory on stack than needed
- ☒ e. The ebp pointers saved on the stack constitute a chain of activation records
- ☒ f. Parameters may be passed in registers or on stack
- ☒ g. Parameters may be passed in registers or on stack
- ☒ h. during execution of a function, ebp is pointing to the old ebp
- ☒ i. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function
- ☒ j. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function
- ☒ k. The return value is either stored on the stack or returned in the eax register

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

## Question 8

Complete

Mark 0.33 out of 0.50

Order the following events in boot process (from 1 onwards)

OS	<input type="text" value="4"/>
Init	<input type="text" value="3"/>
Login interface	<input type="text" value="5"/>
BIOS	<input type="text" value="1"/>
Shell	<input type="text" value="6"/>
Boot loader	<input type="text" value="2"/>

The correct answer is: OS → 3, Init → 4, Login interface → 5, BIOS → 1, Shell → 6, Boot loader → 2

Question **9**

Complete

Mark 0.80 out of 1.00

Match the register with the segment used with it.

edi	ds
esi	ds
esp	ss
ebp	ss
eip	cs

The correct answer is: edi → es, esi → ds, esp → ss, ebp → ss, eip → cs

Question **10**

Complete

Mark 2.00 out of 2.00

Which of the following are NOT a part of job of a typical compiler?

- ☒ a. Check the program for logical errors
- ☒ b. Suggest alternative pieces of code that can be written
- ☐ c. Invoke the linker to link the function calls with their code, extern globals with their declaration
- ☐ d. Convert high level language code to machine code
- ☐ e. Process the # directives in a C program
- ☐ f. Check the program for syntactical errors

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question **11**

Complete

Mark 0.00 out of 2.00

**xv6.img: bootblock kernel**

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- ☒ a. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk.
- ☒ b. The kernel is located at block-1 of the xv6.img
- ☒ c. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)
- ☒ d. The size of the xv6.img is nearly 5 MB
- ☒ e. Blocks in xv6.img after kernel may be all zeroes.
- ☒ f. The bootblock is located on block-0 of the xv6.img
- ☒ g. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk.
- ☒ h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file.
- ☒ i. The bootblock may be 512 bytes or less (looking at the Makefile instruction)
- ☒ j. xv6.img is the virtual processor used by the qemu emulator
- ☒ k. The size of the kernel file is nearly 5 MB

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question **12**

Complete

Mark 0.00 out of 0.50

Is the command "cat README &gt; done &amp;" possible on xv6? (Note the &amp; in the end)

- ☐ a. yes
- ☒ b. no

The correct answer is: yes

Question **13**

Complete

Mark 0.50 out of 0.50

Compare multiprogramming with multitasking

- ☒ a. A multiprogramming system is not necessarily multitasking
- ☐ b. A multitasking system is not necessarily multiprogramming

The correct answer is: A multiprogramming system is not necessarily multitasking

Question **14**

Complete

Mark 0.00 out of 2.00

Select all statements that correctly explain the use/purpose of system calls.

Select one or more:

- ☒ a. Switch from user mode to kernel mode
- ☒ b. Provide services for accessing files
- ☒ c. Allow I/O device access to user processes
- ☒ d. Provide an environment for process creation
- ☒ e. Handle ALL types of interrupts
- ☒ f. Run each instruction of an application program
- ☒ g. Handle exceptions like division by zero

The correct answers are: Switch from user mode to kernel mode, Provide services for accessing files, Allow I/O device access to user processes, Provide an environment for process creation

Question **15**

Complete

Mark 1.00 out of 2.00

Match the program with it's output (ignore newlines in the output. Just focus on the count of the number of 'hi')

```
main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }
```

hi hi

```
main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }
```

hi

```
main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }
```

hi hi

```
main() { int i = NULL; fork(); printf("hi\n"); }
```

hi

The correct answer is: `main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi hi, `main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { int i = NULL; fork(); printf("hi\n"); }` → hi hi

Question **16**

Complete

Mark 1.00 out of  
1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- ☒ a. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- ☒ b. The two modes are essential for a multiprogramming system
- ☒ c. The two modes are essential for a multitasking system
- ☒ d. There is an instruction like 'iret' to return from kernel mode to user mode
- ☒ e. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

[◀ \(Task\) Compulsory xv6 task](#)

Jump to...

[\(Optional Assignment\) Shell Programming\(Conformance tests\) ▶](#)



[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-20-21](#) / [OS-Even-sem-2020-21](#) / 14 March - 20 March  
/ [Quiz - 2 \(18 March\)](#)

**Started on** Thursday, 18 March 2021, 2:46 PM

**State** Finished

**Completed on** Thursday, 18 March 2021, 3:50 PM

**Time taken** 1 hour 4 mins

**Grade** 10.36 out of 20.00 (52%)

Question 1

Partially correct

Mark 0.57 out of 1.00

Mark True, the actions done as part of code of switch() in switch.S, in xv6

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Restore new callee saved registers from kernel stack of new context	✓
<input checked="" type="radio"/>	<input type="radio"/>	Save old callee saved registers on kernel stack of old context	✓
<input type="radio"/>	<input checked="" type="radio"/>	Save old callee saved registers on user stack of old context	✓
<input type="radio"/>	<input checked="" type="radio"/>	Switch from old process context to new process context	✗
<input checked="" type="radio"/>	<input type="radio"/>	Switch from one stack (old) to another(new)	✗
<input type="radio"/>	<input checked="" type="radio"/>	Restore new callee saved registers from user stack of new context	✓
<input type="radio"/>	<input checked="" type="radio"/>	Jump to code in new context	✗

Restore new callee saved registers from kernel stack of new context: True

Save old callee saved registers on kernel stack of old context: True

Save old callee saved registers on user stack of old context: False

Switch from old process context to new process context: False

Switch from one stack (old) to another(new): True

Restore new callee saved registers from user stack of new context: False

Jump to code in new context: False

## Question 2

Partially correct

Mark 0.17 out of 0.50

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 3 entries, right from first line of code of bootloader	✗
kvmalloc() in main()	gdt setup with 5 entries (0 to 4) on one processor	✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
after seginit() in main()	gdt setup with 5 entries (0 to 4) on all processors	✗
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	✗
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

## Question 3

Partially correct

Mark 0.38 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- ☒ a. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- ☒ b. Demand paging requires additional hardware support, compared to paging. ✓
- ☐ c. Paging requires some hardware support in CPU
- ☒ d. With paging, it's possible to have user programs bigger than physical memory. ✗
- ☒ e. Both demand paging and paging support shared memory pages. ✓
- ☐ f. Demand paging always increases effective memory access time.
- ☒ g. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- ☒ h. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- ☐ i. TLB hit ration has zero impact in effective memory access time in demand paging.
- ☐ j. Paging requires NO hardware support in CPU

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

## Question 4

Partially correct

Mark 0.44 out of 0.50

Suppose a processor supports base(relocation register) + limit scheme of MMU.

Assuming this, mark the statements as True/False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The OS may terminate the process while handling the interrupt of memory violation	✓
<input checked="" type="radio"/>	<input type="radio"/>	The hardware detects any memory access beyond the limit value and raises an interrupt	✓
<input type="radio"/>	<input checked="" type="radio"/>	The hardware may terminate the process while handling the interrupt of memory violation	✗
<input checked="" type="radio"/>	<input type="radio"/>	The OS sets up the relocation and limit registers when the process is scheduled	✓
<input checked="" type="radio"/>	<input type="radio"/>	The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;	✓
<input type="radio"/>	<input checked="" type="radio"/>	The process sets up it's own relocation and limit registers when the process is scheduled	✓
<input type="radio"/>	<input checked="" type="radio"/>	The OS detects any memory access beyond the limit value and raises an interrupt	✓
<input type="radio"/>	<input checked="" type="radio"/>	The compiler generates machine code assuming appropriately sized segments for code, data and stack.	✓

The OS may terminate the process while handling the interrupt of memory violation: True

The hardware detects any memory access beyond the limit value and raises an interrupt: True

The hardware may terminate the process while handling the interrupt of memory violation: False

The OS sets up the relocation and limit registers when the process is scheduled: True

The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;: True

The process sets up it's own relocation and limit registers when the process is scheduled: False

The OS detects any memory access beyond the limit value and raises an interrupt: False

The compiler generates machine code assuming appropriately sized segments for code, data and stack.: False

## Question 5

Correct

Mark 0.50 out of 0.50

Consider the following list of free chunks, in continuous memory management:

10k, 25k, 12k, 7k, 9k, 13k

Suppose there is a request for chunk of size 9k, then the free chunk selected under each of the following schemes will be

Best fit:

9k



First fit:

10k



Worst fit:

25k



## Question 6

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about MMU and it's functionality

Select one or more:

- ☐ a. MMU is a separate chip outside the processor
- ☒ b. MMU is inside the processor
- ☐ c. Logical to physical address translations in MMU are done with specific machine instructions
- ☒ d. The operating system interacts with MMU for every single address translation
- ☒ e. Illegal memory access is detected in hardware by MMU and a trap is raised
- ☐ f. The Operating system sets up relevant CPU registers to enable proper MMU translations
- ☒ g. Logical to physical address translations in MMU are done in hardware, automatically
- ☐ h. Illegal memory access is detected by operating system



Your answer is partially correct.

You have correctly selected 3.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

## Question 7

Incorrect

Mark 0.00 out of 0.50

Assuming a 8- KB page size, what is the page numbers for the address 874815 reference in decimal :  
(give answer also in decimal)

Answer: 2186 ✖

The correct answer is: 107

## Question 8

Incorrect

Mark 0.00 out of 0.25

Select the compiler's view of the process's address space, for each of the following MMU schemes:  
(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Segmentation, then paging	Many continuous chunks each of page size	✖
Relocation + Limit	Many continuous chunks of same size	✖
Segmentation	one continuous chunk	✖
Paging	many continuous chunks of variable size	✖

Your answer is incorrect.

The correct answer is: Segmentation, then paging → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Paging → one continuous chunk

## Question 9

Incorrect

Mark 0.00 out of 0.50

Suppose the memory access time is 180ns and TLB hit ratio is 0.3, then effective memory access time is (in nanoseconds);

Answer: 192 ✖

The correct answer is: 306.00

Question **10**

Correct

Mark 0.50 out of 0.50

In xv6, The struct context is given as

```
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Select all the reasons that explain why only these 5 registers are included in the struct context.

- ☒ a. The segment registers are same across all contexts, hence they need not be saved ✓
- ☒ b. esp is not saved in context, because context{} is on stack and it's address is always argument to switch() ✓
- ☐ c. xv6 tries to minimize the size of context to save memory space
- ☐ d. esp is not saved in context, because it's not part of the context
- ☒ e. eax, ecx, edx are caller save, hence no need to save ✓

Your answer is correct.

The correct answers are: The segment registers are same across all contexts, hence they need not be saved, eax, ecx, edx are caller save, hence no need to save, esp is not saved in context, because context{} is on stack and it's address is always argument to switch()

Question **11**

Partially correct

Mark 0.83 out of 1.50

Arrange the following events in order, in page fault handling:

Disk interrupt wakes up the process	7	✓
The reference bit is found to be invalid by MMU	1	✓
OS makes available an empty frame	6	✗
Restart the instruction that caused the page fault	9	✓
A hardware interrupt is issued	3	✗
OS schedules a disk read for the page (from backing store)	5	✓
Process is kept in wait state	4	✗
Page tables are updated for the process	8	✓
Operating system decides that the page was not in memory	2	✗

Your answer is partially correct.

You have correctly selected 5.

The correct answer is: Disk interrupt wakes up the process → 7, The reference bit is found to be invalid by MMU → 1, OS makes available an empty frame → 4, Restart the instruction that caused the page fault → 9, A hardware interrupt is issued → 2, OS schedules a disk read for the page (from backing store) → 5, Process is kept in wait state → 6, Page tables are updated for the process → 8, Operating system decides that the page was not in memory → 3

Question **12**

Incorrect

Mark 0.00 out of 0.50

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

00001010

Now, there is a request for a chunk of 70 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer:



The correct answer is: 11111010

Question **13**

Incorrect

Mark 0.00 out of 0.25

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- ☐ a. end, 4MB
- ☐ b. P2V(end), P2V(PHYSTOP)
- ☐ c. end, P2V(4MB + PHYSTOP)
- ☒ d. P2V(end), PHYSTOP
- ☐ e. end, (4MB + PHYSTOP)
- ☐ f. end, PHYSTOP
- ☐ g. end, P2V(PHYSTOP)



Your answer is incorrect.

The correct answer is: end, P2V(PHYSTOP)



## Question 14

Partially correct

Mark 0.33 out of 0.50

Match the pair

Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓
Inverted Page table	Linear/Parallel search using frame number in page table	✗
Hierarchical Paging	More memory access time per hierarchy	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically, Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy

## Question 15

Partially correct

Mark 0.29 out of 0.50

After virtual memory is implemented

(select T/F for each of the following)One Program's size can be larger than physical memory size

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Code need not be completely in memory	✓
<input checked="" type="radio"/>	<input type="radio"/>	Cumulative size of all programs can be larger than physical memory size	✓
<input type="radio"/>	<input checked="" type="radio"/>	Virtual access to memory is granted	✗
<input checked="" type="radio"/>	<input type="radio"/>	Logical address space could be larger than physical address space	✓
<input type="radio"/>	<input checked="" type="radio"/>	Virtual addresses are available	✗
<input type="radio"/>	<input checked="" type="radio"/>	Relatively less I/O may be possible during process execution	✗
<input checked="" type="radio"/>	<input type="radio"/>	One Program's size can be larger than physical memory size	✓

Code need not be completely in memory: True

Cumulative size of all programs can be larger than physical memory size: True

Virtual access to memory is granted: False

Logical address space could be larger than physical address space: True

Virtual addresses are available: False

Relatively less I/O may be possible during process execution: True

One Program's size can be larger than physical memory size: True

Question **16**

Partially correct

Mark 0.64 out of 1.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only Mark statements True or False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input type="radio"/>	<input checked="" type="radio"/>	The free page-frame are created out of nearly 222 MB	✗
<input checked="" type="radio"/>	<input type="radio"/>	The kernel code and data take up less than 2 MB space	✓
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✗
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓
<input type="radio"/>	<input checked="" type="radio"/>	xv6 uses physical memory upto 224 MB only	✗
<input checked="" type="radio"/>	<input type="radio"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input type="radio"/>	<input checked="" type="radio"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✗

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The free page-frame are created out of nearly 222 MB: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

xv6 uses physical memory upto 224 MB only: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question **17**

Incorrect

Mark 0.00 out of 1.50

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using LRU replacement is:

Answer:



#6# 6,4# 6,4,2 # 0,4,2#0,1,2#6,1,2#6,9,2#0,9,2#0,5,2

The correct answer is: 9

## Question 18

Partially correct

Mark 0.31 out of 0.50

Consider the image given below, which explains how paging works.

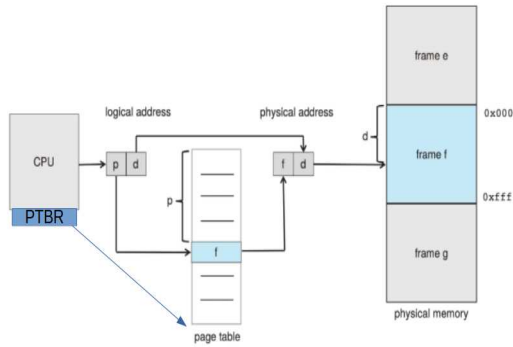


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register	✓
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number	✗
<input checked="" type="radio"/>	<input type="radio"/>	The locating of the page table using PTBR also involves paging translation	✗
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory	✓
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number	✗
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address	✓

The PTBR is present in the CPU as a register: True

The page table is indexed using frame number: False

The page table is indexed using page number: True

The locating of the page table using PTBR also involves paging translation: False

Size of page table is always determined by the size of RAM: False

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

The physical address may not be of the same size (in bits) as the logical address: True

Question **19**

Correct

Mark 2.00 out of 2.00

Given below is shared memory code with two processes sharing a memory segment.

The first process sends a user input string to second process. The second capitalizes the string. Then the first process prints the capitalized version.

Fill in the blanks to complete the code.

**// First process**

#define SHMSZ 27

int main()

```
{
    char c;
    int shmid;
    key_t key;
    char *shm, *s, string[128];
    key = 5679;
    if ((shmid =
```

shmget

```
✓ (key, SHMSZ, IPC_CREAT | 0666)) < 0) {
    perror("shmget");
    exit(1);
}
```

if ((shm =

shmat

```
✓ (shmid, NULL, 0)) == (char *) -1) {
    perror("shmat");
    exit(1);
}
s = shm;
*s = '$';
scanf("%s", string);
strcpy(s + 1, string);
*s = '
```

@

```
✓ ' '; //note the quotes
while(*s != '

```

\$

```
✓ ' )
    sleep(1);
    printf("%s\n", s + 1);
    exit(0);
}
```

**//Second process**

#define SHMSZ 27

int main()

```
{
    int shmid;
    key_t key;
    char *shm, *s;
    int i;
    char string[128];
    key =
```

5679

```

✓ ;
if ((shm = shmget(key, SHMSZ, 0666)) < 0) {
    perror("shmget");
    exit(1);
}
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
    perror("shmat");
    exit(1);
}
s =

```

shm

```

✓ ;
while(*s != '@')
    sleep(1);
for(i = 0; i < strlen(s + 1); i++)
    s[i + 1] = toupper(s[i + 1]);
*s = '$';
exit(0);
}

```

## Question 20

Partially correct

Mark 0.25 out of 0.50

Map the functionality/use with function/variable in xv6 code.

return a free page, if available; 0, otherwise

kinit1()

✗

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

mappages()

✓

Array listing the kernel memory mappings, to be used by setupkvm()

kmap[]

✓

Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

kvmalloc()

✗

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

walkpgdir()

✓

Setup kernel part of a page table, and switch to that page table

setupkvm()

✗

Your answer is partially correct.

You have correctly selected 3.

The correct answer is: return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[], Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Setup kernel part of a page table, and switch to that page table → kvmalloc()

Question **21**

Partially correct

Mark 1.53 out of 2.50

Order events in xv6 timer interrupt code

(Transition from process P1 to P2's code.)

P2 is selected and marked RUNNING	12	✓
Change of stack from user stack to kernel stack of P1	3	✓
Timer interrupt occurs	2	✓
alltraps() will call iret	17	✗
change to context of P2, P2's kernel stack in use now	13	✓
P2's trap() will return to alltraps	16	✗
jump in vector.S	4	✓
P2 will return from sched() in yield()	14	✗
yield() is called	8	✓
trap() is called	7	✓
Process P2 is executing	18	✗
P1 is marked as RUNNABLE	9	✓
P2's yield() will return in trap()	15	✗
Process P1 is executing	1	✓
sched() is called,	11	✗
change to context of the scheduler, scheduler's stack in use now	10	✗
jump to alltraps	5	✓
Trapframe is built on kernel stack of P1	6	✓

Your answer is partially correct.

You have correctly selected 11.

The correct answer is: P2 is selected and marked RUNNING → 12, Change of stack from user stack to kernel stack of P1 → 3, Timer interrupt occurs → 2, alltraps() will call iret → 18, change to context of P2, P2's kernel stack in use now → 13, P2's trap() will return to alltraps → 17, jump in vector.S → 4, P2 will return from sched() in yield() → 15, yield() is called → 8, trap() is called → 7, Process P2 is executing → 14, P1 is marked as RUNNABLE → 9, P2's yield() will return in trap() → 16, Process P1 is executing → 1, sched() is called, → 10, change to context of the scheduler, scheduler's stack in use now → 11, jump to alltraps → 5, Trapframe is built on kernel stack of P1 → 6

Question **22**

Incorrect

Mark 0.00 out of 1.00

Given that the memory access time is 200 ns, probability of a page fault is 0.7 and page fault handling time is 8 ms,  
The effective memory access time in nanoseconds is:

Answer:  ❌

The correct answer is: 5600060.00

Question **23**

Correct

Mark 0.25 out of 0.25

Select the state that is not possible after the given state, for a process:

New:  ✔️

Ready :  ✔️

Running: :  ✔️

Waiting:  ✔️

Question **24**

Partially correct

Mark 0.63 out of 1.00

Select the correct statements about sched() and scheduler() in xv6 code

- ☒ a. scheduler() switches to the selected process's context ✔️
- ☒ b. When either sched() or scheduler() is called, it does not return immediately to caller ✔️
- ☐ c. After call to switch() in sched(), the control moves to code in scheduler()
- ☒ d. Each call to sched() or scheduler() involves change of one stack inside switch() ✔️
- ☐ e. After call to switch() in scheduler(), the control moves to code in sched()
- ☒ f. When either sched() or scheduler() is called, it results in a context switch ✔️
- ☒ g. sched() switches to the scheduler's context ✔️
- ☐ h. sched() and scheduler() are co-routines

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to switch() in scheduler(), the control moves to code in sched(), After call to switch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside switch()



Question **25**

Correct

Mark 0.25 out of 0.25

The data structure used in `kalloc()` and `kfree()` in `xv6` is

- ☐ a. Doubly linked circular list
- ☐ b. Singly linked circular list
- ☐ c. Double linked NULL terminated list
- ☒ d. Singly linked NULL terminated list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

[◀ \(Assignment\) lseek system call in xv6](#)

Jump to...

**Started on** Tuesday, 22 March 2022, 1:55:19 PM

**State** Finished

**Completed on** Tuesday, 22 March 2022, 5:05:18 PM

**Time taken** 3 hours 9 mins

**Grade** 24.68 out of 40.00 (62%)

Question **1**

Correct

Mark 1.00 out of 1.00

Given that a kernel has 1000 KB of total memory, and holes of sizes (in that order) 300 KB, 200 KB, 100 KB, 250 KB. For each of the requests on the left side, match it with the chunk chosen using the specified algorithm.

Consider each request as first request.

100 KB, worst fit	300 KB	✓
150 KB, best fit	200 KB	✓
220 KB, best fit	250 KB	✓
200 KB, first fit	300 KB	✓
150 KB, first fit	300 KB	✓
50 KB, worst fit	300 KB	✓

The correct answer is: 100 KB, worst fit → 300 KB, 150 KB, best fit → 200 KB, 220 KB, best fit → 250 KB, 200 KB, first fit → 300 KB, 150 KB, first fit → 300 KB, 50 KB, worst fit → 300 KB

Question **2**

Incorrect

Mark 0.00 out of 1.00

Given below is a sequence of reference bits on pages before the second chance algorithm runs. Before the algorithm runs, the counter is at the page marked (x). Write the sequence of reference bits after the second chance algorithm has executed once. In the answer write PRECISELY one space BETWEEN each number and do not mention (x).

0 0 1(x) 1 0 1 1

Answer: 1110001



The correct answer is: 0 0 0 0 1 1

## Question 3

Partially correct

Mark 0.50 out of 1.00

Consider a demand-paging system with the following time-measured utilizations:

CPU utilization : 20%

Paging disk: 97.7%

Other I/O devices: 5%

For each of the following, indicate whether it will (or is likely to) improve CPU utilization (even if by a small amount). Explain your answers.

- a. Install a faster CPU :  ✖
- b. Install a bigger paging disk. :  ✔
- c. Increase the degree of multiprogramming. :  ✖
- d. Decrease the degree of multiprogramming. :  ✖
- e. Install more main memory.:  ✔
- f. Install a faster hard disk or multiple controllers with multiple hard disks. :  ✔
- g. Add prepaging to the page-fetch algorithms. :  ✖
- h. Increase the page size. :  ✔

## Question 4

Partially correct

Mark 0.80 out of 1.00

Mark the statements about named and un-named pipes as True or False

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	A named pipe has a name decided by the kernel.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Un-named pipes are inherited by a child process from parent.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Both types of pipes provide FIFO communication.	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Both types of pipes are an extension of the idea of "message passing".	✓
<input type="radio"/>	<input checked="" type="radio"/>	Named pipes can be used for communication between only "related" processes.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.	✓
<input checked="" type="radio"/>	<input checked="" type="radio"/>	The pipe() system call can be used to create either a named or un-named pipe.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Named pipes can exist beyond the life-time of processes using them.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Named pipe exists as a file	✓

A named pipe has a name decided by the kernel.: False

Un-named pipes are inherited by a child process from parent.: True

Both types of pipes provide FIFO communication.: True

The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.: False

Both types of pipes are an extension of the idea of "message passing".: True

Named pipes can be used for communication between only "related" processes.: False

Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.: True

The pipe() system call can be used to create either a named or un-named pipe.: False

Named pipes can exist beyond the life-time of processes using them.: True

Named pipe exists as a file: True

## Question 5

Partially correct

Mark 0.30 out of 1.00

Select the correct statements about interrupt handling in xv6 code

- ☐ a. xv6 uses the 0x64th entry in IDT for system calls
- ☐ b. The function trap() is the called irrespective of hardware interrupt/system-call/exception
- ☒ c. On any interrupt/syscall/exception the control first jumps in trapasm.S ✗
- ☐ d. Before going to alltraps, the kernel stack contains upto 5 entries.
- ☐ e. The CS and EIP are changed only immediately on a hardware interrupt
- ☐ f. The trapframe pointer in struct proc, points to a location on user stack
- ☒ g. All the 256 entries in the IDT are filled ✓
- ☐ h. The function trap() is the called only in case of hardware interrupt
- ☐ i. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- ☒ j. The trapframe pointer in struct proc, points to a location on kernel stack ✓
- ☒ k. xv6 uses the 64th entry in IDT for system calls ✓
- ☒ l. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt ✓
- ☐ m. On any interrupt/syscall/exception the control first jumps in vectors.S

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question **6**

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements w.r.t user and kernel threads

Select one or more:

- ☒ a. A process blocks in many-one model even if a single thread makes a blocking system call ✓
- ☐ b. many-one model can be implemented even if there are no kernel threads
- ☒ c. many-one model gives no speedup on multicore processors ✓
- ☒ d. one-one model increases kernel's scheduling load ✓
- ☐ e. A process may not block in many-one model, if a thread makes a blocking system call
- ☐ f. one-one model can be implemented even if there are no kernel threads
- ☒ g. all three models, that is many-one, one-one, many-many , require a user level thread library ✓

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: many-one model can be implemented even if there are no kernel threads, all three models, that is many-one, one-one, many-many , require a user level thread library, one-one model increases kernel's scheduling load, many-one model gives no speedup on multicore processors, A process blocks in many-one model even if a single thread makes a blocking system call

Question **7**

Correct

Mark 1.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

10011010

Now, there is a request for a chunk of 50 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer:



The correct answer is: 11111010

## Question 8

Incorrect

Mark 0.00 out of 1.00

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 5 entries (0 to 4) on one processor	✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on one processor	✗
kvmalloc() in main()	gdt setup with 3 entries, right from first line of code of bootloader	✗
after seginit() in main()	gdt setup with 5 entries (0 to 4) on all processors	✗
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	✗
entry.S	gdt setup with 5 entries (0 to 4) on all processors	✗

Your answer is incorrect.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

## Question 9

Partially correct

Mark 0.78 out of 1.00

Mark the statements as True or False, w.r.t. mmap()

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	mmap() results in changes to page table of a process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	mmap() is a system call	✓
<input type="radio"/>	<input checked="" type="radio"/>	MAP_SHARED leads to a mapping that is copy-on-write	✓
<input checked="" type="radio"/>	<input type="radio"/>	on failure mmap() returns (void *)-1	✓
<input type="radio"/>	<input checked="" type="radio"/>	on failure mmap() returns NULL	✓
<input type="radio"/>	<input checked="" type="radio"/>	mmap() results in changes to buffer-cache of the kernel.	✓
<input checked="" type="radio"/>	<input type="radio"/>	mmap() can be implemented on both demand paged and non-demand paged systems.	✗
<input type="radio"/>	<input checked="" type="radio"/>	MAP_FIXED guarantees that the mapping is always done at the specified address	✗
<input checked="" type="radio"/>	<input type="radio"/>	MAP_PRIVATE leads to a mapping that is copy-on-write	✓

mmap() results in changes to page table of a process.: True

mmap() is a system call: True

MAP\_SHARED leads to a mapping that is copy-on-write: False

on failure mmap() returns (void \*)-1: True

on failure mmap() returns NULL: False

mmap() results in changes to buffer-cache of the kernel.: False

mmap() can be implemented on both demand paged and non-demand paged systems.: True

MAP\_FIXED guarantees that the mapping is always done at the specified address: False

MAP\_PRIVATE leads to a mapping that is copy-on-write: True



## Question 10

Partially correct

Mark 0.67 out of 1.00




Figure 9.8 Paging hardware.

Mark the statements as True or False, w.r.t. the above diagram (note that the diagram does not cover all details of what actually happens!)

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The logical address issued by CPU is the same one generated by compiler	✗
<input checked="" type="radio"/>	<input type="radio"/>	The combining of f and d is done by MMU	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is in physical memory and must be continuous	✓
<input type="radio"/>	<input checked="" type="radio"/>	Using the offset d in the physical page-frame is done by MMU	✗
<input checked="" type="radio"/>	<input type="radio"/>	The split of logical address into p and d is done by MMU	✓
<input type="radio"/>	<input checked="" type="radio"/>	There are total 3 memory references in this diagram	✓

The logical address issued by CPU is the same one generated by compiler: True

The combining of f and d is done by MMU: True

The page table is in physical memory and must be continuous: True

Using the offset d in the physical page-frame is done by MMU: False

The split of logical address into p and d is done by MMU: True

There are total 3 memory references in this diagram: False

Question **11**

Correct

Mark 1.00 out of 1.00

If one thread opens a file with read privileges then

Select one:

- ☒ a. other threads in the same process can also read from that file
- ☐ b. other threads in the another process can also read from that file
- ☐ c. none of these
- ☐ d. any other thread cannot read from that file



Your answer is correct.

The correct answer is: other threads in the same process can also read from that file

Question **12**

Correct

Mark 2.00 out of 2.00

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using FIFO replacement is:

Answer: 10



#6# 6,4# 6,4,2 #0,4,2# 0,1,2 #0,1,6 #9,1,6# 9,2,6# 9,2,0 #5,2,0

The correct answer is: 10

Question **13**

Correct

Mark 1.00 out of 1.00

For the reference string

3 4 3 5 2

using LRU replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the most correct statement.

Select one:

- ☒ a. LRU will never exhibit Balady's anomaly
- ☐ b. Exhibit Balady's anomaly between 3 and 4 frames
- ☐ c. Exhibit Balady's anomaly between 2 and 3 frames
- ☐ d. This example does not exhibit Balady's anomaly



Your answer is correct.

The correct answer is: LRU will never exhibit Balady's anomaly

## Question 14

Correct

Mark 1.00 out of 1.00

The data structure used in `kalloc()` and `kfree()` in xv6 is

- ☒ a. Singly linked NULL terminated list
- ☐ b. Doubly linked circular list
- ☐ c. Singly linked circular list
- ☐ d. Double linked NULL terminated list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

## Question 15

Partially correct

Mark 0.67 out of 1.00

Map the functionality/use with function/variable in xv6 code.

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

mappages()



Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

kinit1()



Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

walkpgdir()



Setup kernel part of a page table, and switch to that page table

setupkvm()



return a free page, if available; 0, otherwise

kalloc()



Array listing the kernel memory mappings, to be used by `setupkvm()`

kmap[]



Your answer is partially correct.

You have correctly selected 4.

The correct answer is: Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → `mappages()`, Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → `setupkvm()`, Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → `walkpgdir()`, Setup kernel part of a page table, and switch to that page table → `kvmalloc()`, return a free page, if available; 0, otherwise → `kalloc()`, Array listing the kernel memory mappings, to be used by `setupkvm()` → `kmap[]`

## Question 16

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about linking and loading.

Select one or more:


- ☒ a. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- ☒ b. Loader is part of the operating system ✓
- ☒ c. Dynamic linking essentially results in relocatable code. ✓
- ☒ d. Loader is last stage of the linker program ✗
- ☐ e. Continuous memory management schemes can support dynamic linking and dynamic loading.
- ☐ f. Dynamic linking and loading is not possible without demand paging or demand segmentation.
- ☒ g. Static linking leads to non-relocatable code ✗
- ☐ h. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently)
- ☒ i. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!



**Figure 3.2** Diagram of process state.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler	✓
<input type="radio"/>	<input checked="" type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first	✓
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet	✓
<input type="radio"/>	<input checked="" type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.	✓

Only a process in READY state is considered by scheduler: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

Question **18**  
Partially correct  
Mark 0.57 out of 1.00

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	Virtual access to memory is granted to all processes	✓
<input checked="" type="radio"/>	<input type="radio"/>	Virtual addresses become available to executing process	✗
<input checked="" type="radio"/>	<input type="radio"/>	Cumulative size of all programs can be larger than physical memory size	✓
<input checked="" type="radio"/>	<input type="radio"/>	Relatively less I/O may be possible during process execution	✓
<input type="radio"/>	<input checked="" type="radio"/>	Logical address space could be larger than physical address space	✗
<input type="radio"/>	<input checked="" type="radio"/>	One Program's size can be larger than physical memory size	✗
<input checked="" type="radio"/>	<input type="radio"/>	Code need not be completely in memory	✓

Virtual access to memory is granted to all processes: False

Virtual addresses become available to executing process: False

Cumulative size of all programs can be larger than physical memory size: True

Relatively less I/O may be possible during process execution: True

Logical address space could be larger than physical address space: True

One Program's size can be larger than physical memory size: True

Code need not be completely in memory: True

Question **19**  
Correct  
Mark 1.00 out of 1.00

Consider a computer system with a 32-bit logical address and 4- KB page size. The system supports up to 512 MB of physical memory. How many entries are there in each of the following?

Write answer as a decimal number.

A conventional, single-level page table: 1048576 ✓

An inverted page table: 131072 ✓

Question **20**  
Partially correct  
Mark 0.50 out of 1.00

Mark the statements as True or False, w.r.t. thrashing

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The working set model is an attempt at approximating the locality of a process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.	✓
<input checked="" type="radio"/>	<input type="radio"/>	During thrashing the CPU is under-utilised as most time is spent in I/O	✓
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing can be limited if local replacement is used.	✗
<input type="radio"/>	<input checked="" type="radio"/>	mmap() solves the problem of thrashing.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing occurs when the total size of all processes's locality exceeds total memory size.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.	✗
<input type="radio"/>	<input checked="" type="radio"/>	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.	✗
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing can occur even if entire memory is not in use.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing occurs because some process is doing lot of disk I/O.	✗

The working set model is an attempt at approximating the locality of a process.: True

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Thrashing can be limited if local replacement is used.: True

mmap() solves the problem of thrashing.: False

Thrashing occurs when the total size of all processes's locality exceeds total memory size.: True

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

Thrashing can occur even if entire memory is not in use.: False

Thrashing occurs because some process is doing lot of disk I/O.: False

Question **21**

Incorrect

Mark 0.00 out of  
1.00

Select all the correct statements about MMU and it's functionality (on a non-demand paged system)

Select one or more:

- ☐ a. The Operating system sets up relevant CPU registers to enable proper MMU translations
- ☐ b. MMU is a separate chip outside the processor
- ☒ c. Illegal memory access is detected in hardware by MMU and a trap is raised ✓
- ☒ d. Logical to physical address translations in MMU are done with specific machine instructions ✗
- ☒ e. The operating system interacts with MMU for every single address translation ✗
- ☒ f. MMU is inside the processor ✓
- ☒ g. Illegal memory access is detected by operating system ✗
- ☐ h. Logical to physical address translations in MMU are done in hardware, automatically

Your answer is incorrect.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised



Choice of the global or local replacement strategy is a subjective choice for kernel programmers. There are advantages and disadvantages on either side. Out of the following statements, that advocate either global or local replacement strategy, select those statements that have a logically CONSISTENT argument. (That is any statement that is logically correct about either global or local replacement)

Consistent	Inconsistent		
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).	✓
<input type="radio"/>	<input checked="" type="radio"/>	Global replacement may give highly variable per process completion time because number of page faults become unpredictable.	✗
<input type="radio"/>	<input checked="" type="radio"/>	Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.	✗

Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).: Consistent  
 Global replacement may give highly variable per process completion time because number of page faults become unpredictable.: Consistent  
 Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.: Consistent  
 Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.: Consistent  
 Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.: Consistent

Question **23**  
Partially correct  
Mark 0.25 out of 1.00

Select all the correct statements about signals

Select one or more:

- ☐ a. Signal handlers once replaced can't be restored
- ☐ b. The signal handler code runs in kernel mode of CPU
- ☐ c. SIGKILL definitely kills a process because it can't be caught or ignored, and it's default action terminates the process
- ☒ d. SIGKILL definitely kills a process because it's code runs in kernel mode of CPU ✗
- ☒ e. Signals are delivered to a process by another process ✗
- ☒ f. The signal handler code runs in user mode of CPU ✓
- ☒ g. Signals are delivered to a process by kernel ✓
- ☒ h. A signal handler can be invoked asynchronously or synchronously depending on signal type ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Signals are delivered to a process by kernel, A signal handler can be invoked asynchronously or synchronously depending on signal type, The signal handler code runs in user mode of CPU, SIGKILL definitely kills a process because it can't be caught or ignored, and it's default action terminates the process

Question **24**  
Correct  
Mark 1.00 out of 1.00

Select the most common causes of use of IPC by processes

- ☒ a. Breaking up a large task into small tasks and speeding up computation, on multiple core machines ✓
- ☐ b. Get the kernel performance statistics
- ☒ c. Sharing of information of common interest ✓
- ☒ d. More modular code ✓
- ☐ e. More security checks

The correct answers are: Sharing of information of common interest, Breaking up a large task into small tasks and speeding up computation, on multiple core machines, More modular code

Question **25**

Partially correct

Mark 0.10 out of  
2.00

Order the following events, in the creation of init() process in xv6:

1. ☐ sys\_exec runs
2. ☐ code is set to start in forkret() when process gets scheduled
3. ☐ userinit() is called
4. ☐ function pointer from syscalls[] array is invoked
5. ☐ initcode is selected by scheduler for execution
6. ☐ memory mappings are created for "/init" process
7. ☐ kernel stack is allocated for initcode process
8. ☒ values are set in the trapframe of initcode
9. ☐ initcode process is set to be runnable
10. ☐ empty struct proc is obtained for initcode
11. ☐ initcode calls exec system call
12. ☐ the header of "/init" ELF file is ready by kernel
13. ☐ page table mappings of 'initcode' are replaced by mappings of 'init'
14. ☐ initcode process runs
15. ☐ trap() runs
16. ☐ trapframe and context pointers are set to proper location
17. ☐ Stack is allocated for "/init" process
18. ☐ name of process "/init" is copied in struct proc
19. ☐ Arguments on setup on process stack for /init

Your answer is partially correct.

Grading type: Relative to the next item (including last)

Grade details: 1 / 19 = 5%

Here are the scores for each item in this response:

1. 0 / 1 = 0%
2. 0 / 1 = 0%
3. 0 / 1 = 0%
4. 0 / 1 = 0%
5. 0 / 1 = 0%
6. 0 / 1 = 0%
7. 0 / 1 = 0%
8. 1 / 1 = 100%
9. 0 / 1 = 0%
10. 0 / 1 = 0%
11. 0 / 1 = 0%
12. 0 / 1 = 0%
13. 0 / 1 = 0%
14. 0 / 1 = 0%
15. 0 / 1 = 0%
16. 0 / 1 = 0%
17. 0 / 1 = 0%
18. 0 / 1 = 0%
19. 0 / 1 = 0%

The correct order for these items is as follows:

1. userinit() is called
2. empty struct proc is obtained for initcode
3. kernel stack is allocated for initcode process
4. trapframe and context pointers are set to proper location
5. code is set to start in forkret() when process gets scheduled
6. kernel memory mappings are created for initcode
7. values are set in the trapframe of initcode
8. initcode process is set to be runnable
9. initcode is selected by scheduler for execution
10. initcode process runs
11. initcode calls exec system call
12. trap() runs
13. function pointer from syscalls[] array is invoked
14. sys\_exec runs
15. the header of "/init" ELF file is ready by kernel
16. memory mappings are created for "/init" process
17. Stack is allocated for "/init" process
18. Arguments on setup on process stack for /init
19. name of process "/init" is copied in struct proc
20. page table mappings of 'initcode' are replaced by mappings of 'init'

Question **26**

Partially correct

Mark 1.56 out of 2.00

Match the description of a memory management function with the name of the function that provides it, in xv6

Load contents from ELF into pages after allocating the pages first	inituvm()	✗
Switch to user page table	setupkvm()	✗
setup the kernel part in the page table	setupkvm()	✓
Load contents from ELF into existing pages	loaduvm()	✓
Create a copy of the page table of a process	copyuvm()	✓
Copy the code pages of a process	No such function	✓
Mark the page as in-accessible	clearpteu()	✓
Setup and load the user page table for initcode process	inituvm()	✓
Switch to kernel page table	switchkvm()	✓

The correct answer is: Load contents from ELF into pages after allocating the pages first → No such function, Switch to user page table → switchuvm(), setup the kernel part in the page table → setupkvm(), Load contents from ELF into existing pages → loaduvm(), Create a copy of the page table of a process → copyuvm(), Copy the code pages of a process → No such function, Mark the page as in-accessible → clearpteu(), Setup and load the user page table for initcode process → inituvm(), Switch to kernel page table → switchkvm()

## Question 27

Partially correct

Mark 0.75 out of 1.00

Select the correct points of comparison between POSIX and System V shared memory.

- ☐ a. System V is more prevalent than POSIX even today
- ☒ b. POSIX allows giving name to shared memory, System V does not
- ☒ c. POSIX shared memory is newer than System V shared memory
- ☒ d. POSIX shared memory is "thread safe", System V is not



The correct answers are: POSIX shared memory is newer than System V shared memory, POSIX shared memory is "thread safe", System V is not, POSIX allows giving name to shared memory, System V does not, System V is more prevalent than POSIX even today

## Question 28

Incorrect

Mark 0.00 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1, user code executing  
 Timer interrupt  
 Context changes to kernel context  
 Generic interrupt handler runs  
 Generic interrupt handler calls Scheduler  
 Scheduler selects P2 for execution  
 After scheduler, Process P2 user code executing

This sequence of events is:  ✗

Because

✗

## Question 29

Partially correct

Mark 0.67 out of 1.00

W.r.t. xv6 code, match the state of a process with a code that sets the state

UNUSED	<input type="text" value="wait() called by the exiting process itself"/>	✗
SLEEPING	<input type="text" value="sleep(), called by any process blocking itself"/>	✓
RUNNING	<input type="text" value="scheduler()"/>	✓
EMBRYO	<input type="text" value="fork()-&gt;allocproc() before setting up the UVM"/>	✓
RUNNABLE	<input type="text" value="wakeup(), called by an interrupt handler"/>	✓
ZOMBIE	<input type="text" value="exit(), called by an interrupt handler"/>	✗

The correct answer is: UNUSED → wait(), called by parent process, SLEEPING → sleep(), called by any process blocking itself, RUNNING → scheduler(), EMBRYO → fork()->allocproc() before setting up the UVM, RUNNABLE → wakeup(), called by an interrupt handler, ZOMBIE → exit(), called by process itself

## Question 30

Partially correct

Mark 0.33 out of 1.00

Select all correct statements w.r.t. Major and Minor page faults on Linux

- ☐ a. Minor page fault may occur because of a page fault during fork(), on code of an already running process
- ☒ b. Thrashing is possible only due to major page faults ✓
- ☐ c. Major page faults are likely to occur in more numbers at the beginning of the process
- ☒ d. Minor page fault may occur because the page was a shared memory page ✓
- ☐ e. Minor page fault may occur because the page was freed, but still tagged and available in the free page list
- ☐ f. Minor page faults are an improvement of the page buffering techniques

The correct answers are: Minor page fault may occur because the page was a shared memory page, Minor page fault may occur because of a page fault during fork(), on code of an already running process, Minor page fault may occur because the page was freed, but still tagged and available in the free page list, Major page faults are likely to occur in more numbers at the beginning of the process, Thrashing is possible only due to major page faults, Minor page faults are an improvement of the page buffering techniques

## Question 31

Correct

Mark 2.00 out of 2.00

For the reference string

3 4 3 5 2

using FIFO replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the correct statement.

Select one:

- ☐ a. Exhibit Balady's anomaly between 3 and 4 frames
- ☒ b. Do not exhibit Balady's anomaly ✓
- ☐ c. Exhibit Balady's anomaly between 2 and 3 frames

Your answer is correct.

The correct answer is: Do not exhibit Balady's anomaly

## Question 32

Correct

Mark 1.00 out of 1.00

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- ☐ a. end, P2V(4MB + PHYSTOP)
- ☐ b. end, (4MB + PHYSTOP)
- ☐ c. end, PHYSTOP
- ☒ d. end, P2V(PHYSTOP) ✓
- ☐ e. P2V(end), PHYSTOP
- ☐ f. end, 4MB
- ☐ g. P2V(end), P2V(PHYSTOP)

Your answer is correct.

The correct answer is: end, P2V(PHYSTOP)

Question **33**  
 Partially correct  
 Mark 1.09 out of 2.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only Mark statements True or False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✗
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✗
<input checked="" type="radio"/>	<input type="radio"/>	xv6 uses physical memory upto 224 MB only	✓
<input type="radio"/>	<input checked="" type="radio"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✗
<input checked="" type="radio"/>	<input type="radio"/>	The kernel code and data take up less than 2 MB space	✗
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input checked="" type="radio"/>	<input type="radio"/>	The free page-frame are created out of nearly 222 MB	✗
<input checked="" type="radio"/>	<input type="radio"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True  
 The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False  
 The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False  
 xv6 uses physical memory upto 224 MB only: True  
 The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False  
 The kernel code and data take up less than 2 MB space: True  
 The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True  
 The stack allocated in entry.S is used as stack for scheduler's context for first processor: True  
 The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True  
 The free page-frame are created out of nearly 222 MB: True

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

Question **34**

Correct

Mark 1.00 out of 1.00

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The functions like <code>argint()</code> , <code>argstr()</code> make the system call arguments available in the kernel.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The arguments are accessed in the kernel code using <code>esp</code> on the trapframe.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Integer arguments are copied from user memory to kernel memory using <code>argint()</code>	✓
<input type="radio"/>	<input checked="" type="radio"/>	Integer arguments are stored in <code>eax</code> , <code>ebx</code> , <code>ecx</code> , etc. registers	✓
<input checked="" type="radio"/>	<input type="radio"/>	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer	✓
<input type="radio"/>	<input checked="" type="radio"/>	The arguments to system call are copied to kernel stack in <code>trapasm.S</code>	✓
<input checked="" type="radio"/>	<input type="radio"/>	The arguments to system call originally reside on process stack.	✓
<input type="radio"/>	<input checked="" type="radio"/>	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.	✓

The functions like `argint()`, `argstr()` make the system call arguments available in the kernel.: True

The arguments are accessed in the kernel code using `esp` on the trapframe.: True

Integer arguments are copied from user memory to kernel memory using `argint()`: True

Integer arguments are stored in `eax`, `ebx`, `ecx`, etc. registers: False

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

The arguments to system call are copied to kernel stack in `trapasm.S`: False

The arguments to system call originally reside on process stack.: True

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False



Question **35**

Partially correct

Mark 0.15 out of 1.00

Select all the correct statements about process states.

Note that in this question you lose marks for every incorrect choice that you make, proportional to actual number of incorrect choices.

- ☒ a. A process becomes ZOMBIE when another process bites into it's memory ✗
- ☐ b. Process state is stored in the processor
- ☐ c. Process state is implemented as a string
- ☒ d. The scheduler can change state of a process from RUNNALBE to RUNNING ✓
- ☒ e. Process state is stored in the PCB ✓
- ☐ f. Process state can be implemented as just a number
- ☒ g. A process becomes ZOMBIE when it calls exit() ✓
- ☐ h. Process state is changed only by interrupt handlers ✗ -> no exit() also changes it, and so does fork() during it's execution
- ☒ i. The scheduler can change state of a process from RUNNALBE to RUNNING and vice-versa ✗

Your answer is partially correct.

You have selected too many options.

The correct answers are: Process state is stored in the PCB, Process state can be implemented as just a number, The scheduler can change state of a process from RUNNALBE to RUNNING, A process becomes ZOMBIE when it calls exit()

[◀ \(Optional Assignment\) lseek system call in xv6](#)

Jump to...

[Feedback on Quiz-2 ▶](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [24 January - 30 January](#)  
/ [Topic-wise Quiz: 1 \(system calls, x86, calling convention\)](#)

**Started on** Monday, 24 January 2022, 8:42:38 PM

**State** Finished

**Completed on** Monday, 24 January 2022, 9:52:11 PM

**Time taken** 1 hour 9 mins

**Grade** 9.63 out of 20.00 (48%)

Question **1**

Complete

Mark 0.00 out of  
2.00

Which of the following instructions should be privileged?

Select one or more:

- ☒ a. Access memory management unit of the processor
- ☒ b. Turn off interrupts.
- ☒ c. Set value of a memory location
- ☒ d. Switch from user to kernel mode.
- ☒ e. Read the clock.
- ☒ f. Modify entries in device-status table
- ☒ g. Access a general purpose register
- ☒ h. Access I/O device.
- ☒ i. Set value of timer.

The correct answers are: Set value of timer., Access memory management unit of the processor, Turn off interrupts., Modify entries in device-status table, Access I/O device., Switch from user to kernel mode.

Question **2**

Complete

Mark 1.00 out of 1.00

Rank the following storage systems from slowest (first) to fastest(last)

You can drag and drop the items below/above each other.

Magnetic tapes

Optical disk

Hard-disk drives

Nonvolatile memory

Main memory

Cache

Registers

Question **3**

Complete

Mark 1.00 out of 1.00

```
int value = 5;
int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("%d", value); /* LINE A */
    }
    return 0;
}
```

What's the value printed here at LINE A?

Answer:

The correct answer is: 5

Question **4**

Complete

Mark 0.00 out of 0.50

Is the terminal a part of the kernel on GNU/Linux systems?

- ☐ a. no
- ☒ b. yes

The correct answer is: no

Question **5**

Complete

Mark 1.00 out of 1.00

Why should a program exist in memory before it starts executing ?

- ☐ a. Because the variables of the program are stored in memory
- ☐ b. Because the memory is volatile
- ☐ c. Because the hard disk is a slow medium
- ☒ d. Because the processor can run instructions and access data only from memory

The correct answer is: Because the processor can run instructions and access data only from memory

Question **6**

Complete

Mark 1.00 out of 1.00

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- ☒ a. It prohibits a user mode process from running privileged instructions
- ☐ b. It prohibits invocation of kernel code completely, if a user program is running
- ☐ c. It disallows hardware interrupts when a process is running
- ☐ d. It prohibits one process from accessing other process's memory

The correct answer is: It prohibits a user mode process from running privileged instructions

## Question 7

Complete

Mark 0.00 out of 2.00

Select all the correct statements about calling convention on x86 32-bit.

- ☒ a. Return address is one location above the ebp
- ☒ b. Parameters are pushed on the stack in left-right order
- ☒ c. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables
- ☒ d. Compiler may allocate more memory on stack than needed
- ☒ e. The ebp pointers saved on the stack constitute a chain of activation records
- ☒ f. Parameters may be passed in registers or on stack
- ☒ g. Parameters may be passed in registers or on stack
- ☒ h. during execution of a function, ebp is pointing to the old ebp
- ☒ i. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function
- ☒ j. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function
- ☒ k. The return value is either stored on the stack or returned in the eax register

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

## Question 8

Complete

Mark 0.33 out of 0.50

Order the following events in boot process (from 1 onwards)

OS	4
Init	3
Login interface	5
BIOS	1
Shell	6
Boot loader	2

The correct answer is: OS → 3, Init → 4, Login interface → 5, BIOS → 1, Shell → 6, Boot loader → 2

Question **9**

Complete

Mark 0.80 out of 1.00

Match the register with the segment used with it.

edi	ds
esi	ds
esp	ss
ebp	ss
eip	cs

The correct answer is: edi → es, esi → ds, esp → ss, ebp → ss, eip → cs

Question **10**

Complete

Mark 2.00 out of 2.00

Which of the following are NOT a part of job of a typical compiler?

- ☒ a. Check the program for logical errors
- ☒ b. Suggest alternative pieces of code that can be written
- ☐ c. Invoke the linker to link the function calls with their code, extern globals with their declaration
- ☐ d. Convert high level language code to machine code
- ☐ e. Process the # directives in a C program
- ☐ f. Check the program for syntactical errors

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question **11**

Complete

Mark 0.00 out of 2.00

**xv6.img: bootblock kernel**

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- ☒ a. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk.
- ☒ b. The kernel is located at block-1 of the xv6.img
- ☒ c. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)
- ☒ d. The size of the xv6.img is nearly 5 MB
- ☒ e. Blocks in xv6.img after kernel may be all zeroes.
- ☒ f. The bootblock is located on block-0 of the xv6.img
- ☒ g. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk.
- ☒ h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file.
- ☒ i. The bootblock may be 512 bytes or less (looking at the Makefile instruction)
- ☒ j. xv6.img is the virtual processor used by the qemu emulator
- ☒ k. The size of the kernel file is nearly 5 MB

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question **12**

Complete

Mark 0.00 out of 0.50

Is the command "cat README &gt; done &amp;" possible on xv6? (Note the &amp; in the end)

- ☐ a. yes
- ☒ b. no

The correct answer is: yes

Question **13**

Complete

Mark 0.50 out of 0.50

Compare multiprogramming with multitasking

- ☒ a. A multiprogramming system is not necessarily multitasking
- ☐ b. A multitasking system is not necessarily multiprogramming

The correct answer is: A multiprogramming system is not necessarily multitasking

Question **14**

Complete

Mark 0.00 out of 2.00

Select all statements that correctly explain the use/purpose of system calls.

Select one or more:

- ☒ a. Switch from user mode to kernel mode
- ☒ b. Provide services for accessing files
- ☒ c. Allow I/O device access to user processes
- ☒ d. Provide an environment for process creation
- ☒ e. Handle ALL types of interrupts
- ☒ f. Run each instruction of an application program
- ☒ g. Handle exceptions like division by zero

The correct answers are: Switch from user mode to kernel mode, Provide services for accessing files, Allow I/O device access to user processes, Provide an environment for process creation

Question **15**

Complete

Mark 1.00 out of 2.00

Match the program with it's output (ignore newlines in the output. Just focus on the count of the number of 'hi')

```
main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }
```

hi hi

```
main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }
```

hi

```
main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }
```

hi hi

```
main() { int i = NULL; fork(); printf("hi\n"); }
```

hi

The correct answer is: `main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi hi, `main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { int i = NULL; fork(); printf("hi\n"); }` → hi hi



Question **16**

Complete

Mark 1.00 out of 1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- ☒ a. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- ☒ b. The two modes are essential for a multiprogramming system
- ☒ c. The two modes are essential for a multitasking system
- ☒ d. There is an instruction like 'iret' to return from kernel mode to user mode
- ☒ e. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

[◀ \(Task\) Compulsory xv6 task](#)

Jump to...

[\(Optional Assignment\) Shell Programming\(Conformance tests\) ▶](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [7 February - 13 February](#)  
/ [Topic-wise Quiz: 2: 9 Feb \(bootloader, memory management basics, x86\)](#)

**Started on** Wednesday, 9 February 2022, 7:00:20 PM

**State** Finished

**Completed on** Wednesday, 9 February 2022, 7:58:44 PM

**Time taken** 58 mins 24 secs

**Grade** 9.00 out of 11.00 (82%)

Question **1**

Complete

Mark 1.00 out of 1.00

ELF Magic number is

- ☐ a. 0xELFELFELF
- ☐ b. 0x464C457FL
- ☐ c. 0
- ☐ d. 0x0x464CELF
- ☐ e. 0xELF
- ☐ f. 0xFFFFFFFF
- ☒ g. 0x464C457FU

The correct answer is: 0x464C457FU

Question **2**

Complete

Mark 1.00 out of 1.00

The right side of line of code "entry = (void (\*)(void))(elf->entry)" means

- ☒ a. Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing
- ☐ b. Get the "entry" in ELF structure and convert it into a function void pointer
- ☐ c. Convert the "entry" in ELF structure into void
- ☐ d. Get the "entry" in ELF structure and convert it into a void pointer

The correct answer is: Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing

Question **3**

Complete

Mark 0.00 out of 1.00

The variable \$stack in entry.S is

- ☐ a. located at 0
- ☐ b. a memory region allocated as a part of entry.S
- ☒ c. located at the value given by %esp as setup by bootmain()
- ☐ d. located at less than 0x7c00
- ☐ e. located at 0x7c00

The correct answer is: a memory region allocated as a part of entry.S

Question **4**

Complete

Mark 1.00 out of 1.00

x86 provides which of the following type of memory management options?

- ☐ a. segmentation only
- ☐ b. segmentation and one level paging
- ☐ c. segmentation and two level paging
- ☐ d. segmentation or one or two level paging
- ☒ e. segmentation and one or two level paging
- ☐ f. segmentation or paging

The correct answer is: segmentation and one or two level paging

Question **5**

Complete

Mark 1.00 out of 1.00

The kernel is loaded at Physical Address

- ☐ a. 0x0010000
- ☒ b. 0x00100000
- ☐ c. 0x80100000
- ☐ d. 0x80000000

The correct answer is: 0x00100000

Question **6**

Not answered

Marked out of 0.50

Match the pairs of which action is taken by whom

Answer:

The correct answer is: kernel

Question **7**

Complete

Mark 1.00 out of 1.00

Why is the code of entry() in Assembly and not in C?

- ☐ a. There is no particular reason, it could also be in C
- ☒ b. Because it needs to setup paging
- ☐ c. Because the symbol entry() is inside the ELF file
- ☐ d. Because the kernel code must begin in assembly

The correct answer is: Because it needs to setup paging

Question **8**

Complete

Mark 1.00 out of 1.00

The kernel ELF file contains how many Program headers?

- ☐ a. 9
- ☐ b. 4
- ☒ c. 3
- ☐ d. 10
- ☐ e. 2

The correct answer is: 3

Question **9**

Complete

Mark 1.00 out of 1.00

The number of GDT entries setup during boot process of xv6 is

- ☐ a. 255
- ☐ b. 256
- ☐ c. 4
- ☒ d. 3
- ☐ e. 2
- ☐ f. 0

The correct answer is: 3

Question **10**

Complete

Mark 1.00 out of 1.00

The ljmp instruction in general does

- ☐ a. change the CS and EIP to 32 bit mode
- ☐ b. change the CS and EIP to 32 bit mode, and jumps to next line of code
- ☐ c. change the CS and EIP to 32 bit mode, and jumps to kernel code
- ☒ d. change the CS and EIP to 32 bit mode, and jumps to new value of EIP

The correct answer is: change the CS and EIP to 32 bit mode, and jumps to new value of EIP

Question **11**

Not answered

Marked out of  
0.50

code line, MMU setting: Match the line of xv6 code with the MMU setup employed

Answer:

The correct answer is: `inb $0x64,%al`

Question **12**

Complete

Mark 1.00 out of  
1.00

which of the following is not a difference between real mode and protected mode

- ☐ a. in real mode general purpose registers are 16 bit, in protected mode they are 32 bit
- ☐ b. in real mode the segment is multiplied by 16, in protected mode segment is used as index in GDT
- ☐ c. in real mode the addressable memory is less than in protected mode
- ☒ d. in real mode the addressable memory is more than in protected mode
- ☐ e. processor starts in real mode

The correct answer is: in real mode the addressable memory is more than in protected mode

◀ Homework questions: Basics of  
MM, xv6 booting

Jump to...

(Code) Files, redirection, dup,  
(IPC)pipe ▶

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [14 February - 20 February](#) / [Topic-wise Quiz-3 \(processes, trap handling, scheduler\)](#)

**Started on** Monday, 21 February 2022, 7:02:09 PM

**State** Finished

**Completed on** Monday, 21 February 2022, 7:57:10 PM

**Time taken** 55 mins 1 sec

**Grade** 8.73 out of 10.00 (87%)

Question **1**

Complete

Mark 1.00 out of 1.00

Arrange in correct order, the files involved in execution of system call

trapasm.S 3

usys.S 1

vectors.S 2

trap.c 4

The correct answer is: trapasm.S → 3, usys.S → 1, vectors.S → 2, trap.c → 4

Question **2**

Complete

Mark 0.33 out of 0.50

Match the MACRO with it's meaning

KERNLINK 2.1 GB

PHYSTOP 224 MB

KERNBASE 2 GB

The correct answer is: KERNLINK → 2.224 GB, PHYSTOP → 224 MB, KERNBASE → 2 GB

Question **3**

Complete

Mark 0.50 out of 1.00

Which of the following is not a task of the code of `swtch()` function

- ☐ a. Switch stacks
- ☒ b. Change the kernel stack location
- ☐ c. Save the old context
- ☐ d. Jump to next context EIP
- ☐ e. Save the return value of the old context code
- ☐ f. Load the new context

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question **4**

Complete

Mark 0.50 out of 0.50

Match the names of PCB structures with kernel

xv6    struct proc

linux    struct task\_struct

The correct answer is: xv6 → struct proc, linux → struct task\_struct

Question **5**

Complete

Mark 0.00 out of 0.50

Which of the following state transitions are not possible?

- ☐ a. Running -> Waiting
- ☒ b. Ready -> Waiting
- ☒ c. Ready -> Terminated
- ☒ d. Waiting -> Terminated

The correct answers are: Ready -> Terminated, Waiting -> Terminated, Ready -> Waiting

## Question 6

Complete

Mark 0.90 out of 1.00

Match the elements of C program to their place in memory

Global Static variables	Data
Mallocated Memory	Heap
Global variables	Data
Local Static variables	Data
Function code	Code
Arguments	Stack
Local Variables	Stack
Code of main()	Code
#define MACROS	No Memory needed
#include files	No Memory needed

The correct answer is: Global Static variables → Data, Mallocated Memory → Heap, Global variables → Data, Local Static variables → Data, Function code → Code, Arguments → Stack, Local Variables → Stack, Code of main() → Code, #define MACROS → No Memory needed, #include files → No memory needed

## Question 7

Complete

Mark 1.00 out of 1.00

The trapframe, in xv6, is built by the

- ☐ a. hardware, vectors.S, trapasm.S, trap()
- ☐ b. vectors.S, trapasm.S
- ☐ c. hardware, trapasm.S
- ☐ d. hardware, vectors.S
- ☒ e. hardware, vectors.S, trapasm.S

The correct answer is: hardware, vectors.S, trapasm.S

## Question 8

Complete

Mark 1.00 out of 1.00

A process blocks itself means

- ☒ a. The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler
- ☐ b. The kernel code of system call calls scheduler
- ☐ c. The application code calls the scheduler
- ☐ d. The kernel code of an interrupt handler, moves the process to a waiting queue and calls scheduler

The correct answer is: The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler



Question **9**

Complete

Mark 1.00 out of 1.00

The "push 0" in vectors.S is

- ☐ a. A placeholder to match the size of struct trapframe
- ☐ b. To indicate that it's a system call and not a hardware interrupt
- ☒ c. Place for the error number value
- ☐ d. To be filled in as the return value of the system call

The correct answer is: Place for the error number value

Question **10**

Complete

Mark 0.50 out of 0.50

Match the File descriptors to their meaning

- 2 Standard error
- 1 Standard output
- 0 Standard Input

The correct answer is: 2 → Standard error, 1 → Standard output, 0 → Standard Input

Question **11**

Complete

Mark 1.00 out of 1.00

Select the odd one out

- ☐ a. Process stack of running process to kernel stack of running process
- ☐ b. Kernel stack of running process to kernel stack of scheduler
- ☒ c. Kernel stack of new process to kernel stack of scheduler
- ☐ d. Kernel stack of scheduler to kernel stack of new process
- ☐ e. Kernel stack of new process to Process stack of new process

The correct answer is: Kernel stack of new process to kernel stack of scheduler

Question **12**

Complete

Mark 1.00 out of  
1.00

What will be the output of this program

```
int main() {  
    int fd;  
    printf("%d ", open("/etc/passwd", O_RDONLY));  
    close(1);  
    fd = printf("%d ", open("/etc/passwd", O_RDONLY));  
    close(fd);  
    fd = printf("%d ", open("/etc/passwd", O_RDONLY));  
}
```

- ☐ a. 1 1 1
- ☐ b. 3 1 2
- ☐ c. 3 3 3
- ☐ d. 3 4 5
- ☐ e. 2 2 2
- ☒ f. 3 1 1

The correct answer is: 3 1 1

[◀ Some possible course mini projects](#)

Jump to...

[Tell Abhijit your travel schedule! ▶](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [21 February - 27 February](#) / [Topic-wise Quiz-4 \(Virtual Memory\)](#)

**Started on** Saturday, 26 February 2022, 5:18:34 PM

**State** Finished

**Completed on** Saturday, 26 February 2022, 6:33:35 PM

**Time taken** 1 hour 15 mins

**Grade** 6.56 out of 15.00 (44%)

Question **1**

Complete

Mark 1.00 out of 1.00

Assuming a 8- KB page size, what is the page numbers for the address 26583 reference in decimal :  
(give answer also in decimal)

Answer:

The correct answer is: 3

Question **2**

Complete

Mark 0.50 out of 1.00

For the reference string  
3 4 3 5 2  
the number of page faults (including initial ones) using

FIFO replacement and 2 page frames is :

FIFO replacement and 3 page frames is :

Question **3**

Complete

Mark 0.52 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- ☒ a. Demand paging always increases effective memory access time.
- ☒ b. With demand paging, it's possible to have user programs bigger than physical memory.
- ☒ c. Calculations of number of bits for page number and offset are same in paging and demand paging.
- ☒ d. TLB hit ration has zero impact in effective memory access time in demand paging.
- ☐ e. With paging, it's possible to have user programs bigger than physical memory.
- ☐ f. Demand paging requires additional hardware support, compared to paging.
- ☒ g. Paging requires some hardware support in CPU
- ☒ h. The meaning of valid-invalid bit in page table is different in paging and demand-paging.
- ☒ i. Both demand paging and paging support shared memory pages.
- ☐ j. Paging requires NO hardware support in CPU

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question **4**

Complete

Mark 0.00 out of 1.00

Calculate the EAT in NANO-seconds (upto 2 decimal points) w.r.t. a page fault, given

Memory access time = 123 ns

Average page fault service time = 9 ms

Page fault rate = 0.7

Answer:

The correct answer is: 6300036.90

Question **5**

Complete

Mark 1.00 out of 1.00

Given six memory partitions of 300 KB , 600 KB , 350 KB , 200 KB , 750 KB , and 125 KB (in order), how would the first-fit, best-fit, and worst-fit algorithms place processes of size 115 KB and 500 KB (in order)?

first fit 115 KB	300 KB
best fit 500 KB	600 KB
worst fit 500 KB	635 KB
worst fit 115 KB	750 KB
best fit 115 KB	125 KB
first fit 500 KB	600 KB

The correct answer is: first fit 115 KB → 300 KB, best fit 500 KB → 600 KB, worst fit 500 KB → 635 KB, worst fit 115 KB → 750 KB, best fit 115 KB → 125 KB, first fit 500 KB → 600 KB

Question 6

Complete

Mark 0.50 out of 1.00

W.r.t the figure given below, mark the given statements as True or False.




Figure 10.9 Need for page replacement.

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement means chose any of the frame from 2 to 7
<input checked="" type="radio"/>	<input type="radio"/>	The kernel's pages can not used for replacement if kernel is not pageable.
<input type="radio"/>	<input checked="" type="radio"/>	Page 1 of process 1 needs a replacement
<input type="radio"/>	<input checked="" type="radio"/>	Handling this scenario demands two disk I/Os
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement means chose any of the frame from 2 to 7
<input checked="" type="radio"/>	<input type="radio"/>	Kernel occupies two page frames
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement means chose any of the frame from 0 to 7
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement means chose any of the frames 2, 3, 6

Local replacement means chose any of the frame from 2 to 7: False

The kernel's pages can not used for replacement if kernel is not pageable.: True

Page 1 of process 1 needs a replacement: True

Handling this scenario demands two disk I/Os: True

Global replacement means chose any of the frame from 2 to 7: True

Kernel occupies two page frames: True

Global replacement means chose any of the frame from 0 to 7: False

Local replacement means chose any of the frames 2, 3, 6: True

Question **7**

Complete

Mark 0.75 out of 1.00

which of the following, do you think, are valid concerns for making the kernel pageable?

- ☒ a. The kernel must have some dedicated frames for it's own work
- ☒ b. The disk driver and disk interrupt handler should not be pageable
- ☒ c. The page fault handler should not be pageable
- ☐ d. No part of kernel code should be pageable.
- ☐ e. The kernel's own page tables should not be pageable
- ☐ f. No data structure of kernel should be pageable

The correct answers are: The kernel's own page tables should not be pageable, The page fault handler should not be pageable, The kernel must have some dedicated frames for it's own work, The disk driver and disk interrupt handler should not be pageable

Question **8**

Complete

Mark 0.25 out of 0.50

Map the technique with it's feature/problem

static loading	wastage of physical memory
dynamic loading	allocate memory only if needed
static linking	small executable file
dynamic linking	large executable file

The correct answer is: static loading → wastage of physical memory, dynamic loading → allocate memory only if needed, static linking → large executable file, dynamic linking → small executable file

## Question 9

Complete

Mark 0.50 out of 1.00

Given below is the output of the command "ps -eo minflt,majflt,cmd" on a Linux Desktop system. Select the statements that are consistent with the output

```
626729 482768 /usr/lib/firefox/firefox -contentproc -parentBuildID 20220202182137
-prefsLen 9256 -prefMapSize 264738 -appDir /usr/lib/firefox/browser 6094 true rdd
2167 687 /usr/sbin/apache2 -k start
1265185 222 /usr/bin/gnome-shell
102648 111 /usr/sbin/mysqld
9813 0 bash
15497 370 /usr/bin/gedit --gapplication-service
```

- ☒ a. The bash shell is mostly busy doing work within a particular locality
- ☐ b. All of the processes here exhibit some good locality of reference
- ☐ c. Apache web-server has not been doing much work
- ☒ d. Firefox has likely been running for a large amount of time

The correct answers are: Firefox has likely been running for a large amount of time, Apache web-server has not been doing much work, The bash shell is mostly busy doing work within a particular locality, All of the processes here exhibit some good locality of reference

## Question 10

Complete

Mark 1.00 out of 1.00

Page sizes are a power of 2 because

Select one:

- ☐ a. Power of 2 calculations are highly efficient
- ☐ b. operating system calculations happen using power of 2
- ☐ c. MMU only understands numbers that are power of 2
- ☒ d. Certain bits are reserved for offset in logical address. Hence page size =  $2^{(\text{no. of offset bits})}$
- ☐ e. Certain bits are reserved for offset in logical address. Hence page size =  $2^{(32 - \text{no. of offset bits})}$

The correct answer is: Certain bits are reserved for offset in logical address. Hence page size =  $2^{(\text{no. of offset bits})}$



Question **11**

Complete

Mark 0.29 out of 1.00

Suppose two processes share a library between them. The library consists of 5 pages, and these 5 pages are mapped to frames 9, 15, 23, 4, 7 respectively. Process P1 has got 6 pages, first 3 of which consist of process's own code/data and 3 correspond to library's pages 0, 2, 4. Process P2 has got 7 pages, first 3 of which consist of process's own code/data and remaining 4 correspond to library's pages 0, 1, 3, 4. Fill in the blanks for page table entries of P1 and P2.

Page table of P2, Page 3	4
Page table of P2, Page 1	9
Page table of P1, Page 3	7
Page table of P2, Page 0	5
Page table of P2, Page 4	6
Page table of P1, Page 5	2
Page table of P1, Page 4	23

The correct answer is: Page table of P2, Page 3 → 4, Page table of P2, Page 1 → 15, Page table of P1, Page 3 → 9, Page table of P2, Page 0 → 9, Page table of P2, Page 4 → 7, Page table of P1, Page 5 → 7, Page table of P1, Page 4 → 23

Question **12**

Complete

Mark 0.25 out of 1.00

Select all the correct statements, w.r.t. Copy on Write

- ☐ a. use of COW during fork() is useless if exec() is called by the child
- ☒ b. use of COW during fork() is useless if child called exit()
- ☒ c. COW helps us save memory
- ☐ d. If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated
- ☒ e. Fork() used COW technique to improve performance of new process creation.
- ☒ f. Vfork() assumes that there will be no write, but rather exec()

The correct answers are: Fork() used COW technique to improve performance of new process creation., If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated, COW helps us save memory, Vfork() assumes that there will be no write, but rather exec()

Question **13**

Complete

Mark 0.00 out of 1.00

Order the following events, related to page fault handling, in correct order

1. Page fault interrupt is generated
2. Disk read is issued
3. Disk interrupt handler runs
4. MMU detects that a page table entry is marked "invalid"
5. Page fault handler in kernel starts executing
6. Empty frame is found
7. Page table of page faulted process is updated
8. Disk Interrupt occurs
9. Page faulted process is moved to ready-queue
10. Other processes scheduled by scheduler
11. Page faulting process is made to wait in a queue
12. Page fault handler detects that it's a page fault and not illegal memory access

The correct order for these items is as follows:

1. MMU detects that a page table entry is marked "invalid"
2. Page fault interrupt is generated
3. Page fault handler in kernel starts executing
4. Page fault handler detects that it's a page fault and not illegal memory access
5. Empty frame is found
6. Disk read is issued
7. Page faulting process is made to wait in a queue
8. Other processes scheduled by scheduler
9. Disk Interrupt occurs
10. Disk interrupt handler runs
11. Page table of page faulted process is updated
12. Page faulted process is moved to ready-queue

## Question 14

Complete

Mark 0.00 out of 1.00

Given below is the "maps" file for a particular instance of "vim.basic" process.

Mark the given statements as True or False, w.r.t. the contents of the map file.

```
55a43501b000-55a435049000 r--p 00000000 103:05 917529
/usr/bin/vim.basic
55a435049000-55a435248000 r-xp 0002e000 103:05 917529
/usr/bin/vim.basic
55a435248000-55a4352b6000 r--p 0022d000 103:05 917529
/usr/bin/vim.basic
55a4352b7000-55a4352c5000 r--p 0029b000 103:05 917529
/usr/bin/vim.basic
55a4352c5000-55a4352e2000 rw-p 002a9000 103:05 917529
/usr/bin/vim.basic
55a4352e2000-55a4352f0000 rw-p 00000000 00:00 0
55a436bc9000-55a436e5b000 rw-p 00000000 00:00 0
7f275b0a3000-7f275b0a6000 r--p 00000000 103:05 917901
/usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f275b0a6000-7f275b0ad000 r-xp 00003000 103:05 917901
/usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f275b0ad000-7f275b0af000 r--p 0000a000 103:05 917901
/usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f275b0af000-7f275b0b0000 r--p 0000b000 103:05 917901
/usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f275b0b0000-7f275b0b1000 rw-p 0000c000 103:05 917901
/usr/lib/x86_64-linux-gnu/libnss_files-2.31.so
7f275b0b1000-7f275b0b7000 rw-p 00000000 00:00 0
7f275b0b7000-7f275b8f5000 r--p 00000000 103:05 925247
/usr/lib/locale/locale-archive
7f275b8f5000-7f275b8fa000 rw-p 00000000 00:00 0
7f275b8fa000-7f275b8fc000 r--p 00000000 103:05 924216
/usr/lib/x86_64-linux-gnu/libogg.so.0.8.4
7f275b8fc000-7f275b901000 r-xp 00002000 103:05 924216
/usr/lib/x86_64-linux-gnu/libogg.so.0.8.4
7f275b901000-7f275b904000 r--p 00007000 103:05 924216
/usr/lib/x86_64-linux-gnu/libogg.so.0.8.4
7f275b904000-7f275b905000 ---p 0000a000 103:05 924216
/usr/lib/x86_64-linux-gnu/libogg.so.0.8.4
7f275b905000-7f275b906000 r--p 0000a000 103:05 924216
/usr/lib/x86_64-linux-gnu/libogg.so.0.8.4
7f275b906000-7f275b907000 rw-p 0000b000 103:05 924216
/usr/lib/x86_64-linux-gnu/libogg.so.0.8.4
7f275b907000-7f275b90a000 r--p 00000000 103:05 924627
/usr/lib/x86_64-linux-gnu/libvorbis.so.0.4.8
7f275b90a000-7f275b921000 r-xp 00003000 103:05 924627
/usr/lib/x86_64-linux-gnu/libvorbis.so.0.4.8
7f275b921000-7f275b932000 r--p 0001a000 103:05 924627
/usr/lib/x86_64-linux-gnu/libvorbis.so.0.4.8
7f275b932000-7f275b933000 ---p 0002b000 103:05 924627
/usr/lib/x86_64-linux-gnu/libvorbis.so.0.4.8
7f275b933000-7f275b934000 r--p 0002b000 103:05 924627
/usr/lib/x86_64-linux-gnu/libvorbis.so.0.4.8
7f275b934000-7f275b935000 rw-p 0002c000 103:05 924627
/usr/lib/x86_64-linux-gnu/libvorbis.so.0.4.8
```

[heap]

```
7f275b935000-7f275b937000 rw-p 00000000 00:00 0
7f275b937000-7f275b938000 r--p 00000000 103:05 917914
/usr/lib/x86_64-linux-gnu/libutil-2.31.so
7f275b938000-7f275b939000 r-xp 00001000 103:05 917914
/usr/lib/x86_64-linux-gnu/libutil-2.31.so
7f275b939000-7f275b93a000 r--p 00002000 103:05 917914
/usr/lib/x86_64-linux-gnu/libutil-2.31.so
7f275b93a000-7f275b93b000 r--p 00002000 103:05 917914
/usr/lib/x86_64-linux-gnu/libutil-2.31.so
7f275b93b000-7f275b93c000 rw-p 00003000 103:05 917914
/usr/lib/x86_64-linux-gnu/libutil-2.31.so
7f275b93c000-7f275b93e000 r--p 00000000 103:05 915906
/usr/lib/x86_64-linux-gnu/libz.so.1.2.11
7f275b93e000-7f275b94f000 r-xp 00002000 103:05 915906
/usr/lib/x86_64-linux-gnu/libz.so.1.2.11
7f275b94f000-7f275b955000 r--p 00013000 103:05 915906
/usr/lib/x86_64-linux-gnu/libz.so.1.2.11
7f275b955000-7f275b956000 ---p 00019000 103:05 915906
/usr/lib/x86_64-linux-gnu/libz.so.1.2.11
7f275b956000-7f275b957000 r--p 00019000 103:05 915906
/usr/lib/x86_64-linux-gnu/libz.so.1.2.11
7f275b957000-7f275b958000 rw-p 0001a000 103:05 915906
/usr/lib/x86_64-linux-gnu/libz.so.1.2.11
7f275b958000-7f275b95c000 r--p 00000000 103:05 923645
/usr/lib/x86_64-linux-gnu/libexpat.so.1.6.11
7f275b95c000-7f275b978000 r-xp 00004000 103:05 923645
/usr/lib/x86_64-linux-gnu/libexpat.so.1.6.11
7f275b978000-7f275b982000 r--p 00020000 103:05 923645
/usr/lib/x86_64-linux-gnu/libexpat.so.1.6.11
7f275b982000-7f275b983000 ---p 0002a000 103:05 923645
/usr/lib/x86_64-linux-gnu/libexpat.so.1.6.11
7f275b983000-7f275b985000 r--p 0002a000 103:05 923645
/usr/lib/x86_64-linux-gnu/libexpat.so.1.6.11
7f275b985000-7f275b986000 rw-p 0002c000 103:05 923645
/usr/lib/x86_64-linux-gnu/libexpat.so.1.6.11
7f275b986000-7f275b988000 r--p 00000000 103:05 924057
/usr/lib/x86_64-linux-gnu/libltdl.so.7.3.1
7f275b988000-7f275b98d000 r-xp 00002000 103:05 924057
/usr/lib/x86_64-linux-gnu/libltdl.so.7.3.1
7f275b98d000-7f275b98f000 r--p 00007000 103:05 924057
/usr/lib/x86_64-linux-gnu/libltdl.so.7.3.1
7f275b98f000-7f275b990000 r--p 00008000 103:05 924057
/usr/lib/x86_64-linux-gnu/libltdl.so.7.3.1
7f275b990000-7f275b991000 rw-p 00009000 103:05 924057
/usr/lib/x86_64-linux-gnu/libltdl.so.7.3.1
7f275b991000-7f275b995000 r--p 00000000 103:05 921934
/usr/lib/x86_64-linux-gnu/libtdb.so.1.4.3
7f275b995000-7f275b9a3000 r-xp 00004000 103:05 921934
/usr/lib/x86_64-linux-gnu/libtdb.so.1.4.3
7f275b9a3000-7f275b9a9000 r--p 00012000 103:05 921934
/usr/lib/x86_64-linux-gnu/libtdb.so.1.4.3
7f275b9a9000-7f275b9aa000 r--p 00017000 103:05 921934
/usr/lib/x86_64-linux-gnu/libtdb.so.1.4.3
```

7f275b9aa000-7f275b9ab000 rw-p 00018000 103:05 921934  
/usr/lib/x86\_64-linux-gnu/libtdb.so.1.4.3  
7f275b9ab000-7f275b9ad000 rw-p 00000000 00:00 0  
7f275b9ad000-7f275b9af000 r--p 00000000 103:05 924631  
/usr/lib/x86\_64-linux-gnu/libvorbisfile.so.3.3.7  
7f275b9af000-7f275b9b4000 r-xp 00002000 103:05 924631  
/usr/lib/x86\_64-linux-gnu/libvorbisfile.so.3.3.7  
7f275b9b4000-7f275b9b5000 r--p 00007000 103:05 924631  
/usr/lib/x86\_64-linux-gnu/libvorbisfile.so.3.3.7  
7f275b9b5000-7f275b9b6000 ---p 00008000 103:05 924631  
/usr/lib/x86\_64-linux-gnu/libvorbisfile.so.3.3.7  
7f275b9b6000-7f275b9b7000 r--p 00008000 103:05 924631  
/usr/lib/x86\_64-linux-gnu/libvorbisfile.so.3.3.7  
7f275b9b7000-7f275b9b8000 rw-p 00009000 103:05 924631  
/usr/lib/x86\_64-linux-gnu/libvorbisfile.so.3.3.7  
7f275b9b8000-7f275b9ba000 r--p 00000000 103:05 924277  
/usr/lib/x86\_64-linux-gnu/libpcre2-8.so.0.9.0  
7f275b9ba000-7f275ba1e000 r-xp 00002000 103:05 924277  
/usr/lib/x86\_64-linux-gnu/libpcre2-8.so.0.9.0  
7f275ba1e000-7f275ba46000 r--p 00066000 103:05 924277  
/usr/lib/x86\_64-linux-gnu/libpcre2-8.so.0.9.0  
7f275ba46000-7f275ba47000 r--p 0008d000 103:05 924277  
/usr/lib/x86\_64-linux-gnu/libpcre2-8.so.0.9.0  
7f275ba47000-7f275ba48000 rw-p 0008e000 103:05 924277  
/usr/lib/x86\_64-linux-gnu/libpcre2-8.so.0.9.0  
7f275ba48000-7f275ba6d000 r--p 00000000 103:05 917893  
/usr/lib/x86\_64-linux-gnu/libc-2.31.so  
7f275ba6d000-7f275bbe5000 r-xp 00025000 103:05 917893  
/usr/lib/x86\_64-linux-gnu/libc-2.31.so  
7f275bbe5000-7f275bc2f000 r--p 0019d000 103:05 917893  
/usr/lib/x86\_64-linux-gnu/libc-2.31.so  
7f275bc2f000-7f275bc30000 ---p 001e7000 103:05 917893  
/usr/lib/x86\_64-linux-gnu/libc-2.31.so  
7f275bc30000-7f275bc33000 r--p 001e7000 103:05 917893  
/usr/lib/x86\_64-linux-gnu/libc-2.31.so  
7f275bc33000-7f275bc36000 rw-p 001ea000 103:05 917893  
/usr/lib/x86\_64-linux-gnu/libc-2.31.so  
7f275bc36000-7f275bc3a000 rw-p 00000000 00:00 0  
7f275bc3a000-7f275bc41000 r--p 00000000 103:05 917906  
/usr/lib/x86\_64-linux-gnu/libpthread-2.31.so  
7f275bc41000-7f275bc52000 r-xp 00007000 103:05 917906  
/usr/lib/x86\_64-linux-gnu/libpthread-2.31.so  
7f275bc52000-7f275bc57000 r--p 00018000 103:05 917906  
/usr/lib/x86\_64-linux-gnu/libpthread-2.31.so  
7f275bc57000-7f275bc58000 r--p 0001c000 103:05 917906  
/usr/lib/x86\_64-linux-gnu/libpthread-2.31.so  
7f275bc58000-7f275bc59000 rw-p 0001d000 103:05 917906  
/usr/lib/x86\_64-linux-gnu/libpthread-2.31.so  
7f275bc59000-7f275bc5d000 rw-p 00000000 00:00 0  
7f275bc5d000-7f275bcce000 r--p 00000000 103:05 917016  
/usr/lib/x86\_64-linux-gnu/libpython3.8.so.1.0  
7f275bcce000-7f275bf29000 r-xp 00071000 103:05 917016  
/usr/lib/x86\_64-linux-gnu/libpython3.8.so.1.0

7f275bf29000-7f275c142000 r--p 002cc000 103:05 917016  
/usr/lib/x86\_64-linux-gnu/libpython3.8.so.1.0  
7f275c142000-7f275c143000 ---p 004e5000 103:05 917016  
/usr/lib/x86\_64-linux-gnu/libpython3.8.so.1.0  
7f275c143000-7f275c149000 r--p 004e5000 103:05 917016  
/usr/lib/x86\_64-linux-gnu/libpython3.8.so.1.0  
7f275c149000-7f275c190000 rw-p 004eb000 103:05 917016  
/usr/lib/x86\_64-linux-gnu/libpython3.8.so.1.0  
7f275c190000-7f275c1b3000 rw-p 00000000 00:00 0  
7f275c1b3000-7f275c1b4000 r--p 00000000 103:05 917894  
/usr/lib/x86\_64-linux-gnu/libdl-2.31.so  
7f275c1b4000-7f275c1b6000 r-xp 00001000 103:05 917894  
/usr/lib/x86\_64-linux-gnu/libdl-2.31.so  
7f275c1b6000-7f275c1b7000 r--p 00003000 103:05 917894  
/usr/lib/x86\_64-linux-gnu/libdl-2.31.so  
7f275c1b7000-7f275c1b8000 r--p 00003000 103:05 917894  
/usr/lib/x86\_64-linux-gnu/libdl-2.31.so  
7f275c1b8000-7f275c1b9000 rw-p 00004000 103:05 917894  
/usr/lib/x86\_64-linux-gnu/libdl-2.31.so  
7f275c1b9000-7f275c1bb000 rw-p 00000000 00:00 0  
7f275c1bb000-7f275c1c0000 r-xp 00000000 103:05 923815  
/usr/lib/x86\_64-linux-gnu/libgpm.so.2  
7f275c1c0000-7f275c3bf000 ---p 00005000 103:05 923815  
/usr/lib/x86\_64-linux-gnu/libgpm.so.2  
7f275c3bf000-7f275c3c0000 r--p 00004000 103:05 923815  
/usr/lib/x86\_64-linux-gnu/libgpm.so.2  
7f275c3c0000-7f275c3c1000 rw-p 00005000 103:05 923815  
/usr/lib/x86\_64-linux-gnu/libgpm.so.2  
7f275c3c1000-7f275c3c3000 r--p 00000000 103:05 923315  
/usr/lib/x86\_64-linux-gnu/libacl.so.1.1.2253  
7f275c3c3000-7f275c3c8000 r-xp 00002000 103:05 923315  
/usr/lib/x86\_64-linux-gnu/libacl.so.1.1.2253  
7f275c3c8000-7f275c3ca000 r--p 00007000 103:05 923315  
/usr/lib/x86\_64-linux-gnu/libacl.so.1.1.2253  
7f275c3ca000-7f275c3cb000 r--p 00008000 103:05 923315  
/usr/lib/x86\_64-linux-gnu/libacl.so.1.1.2253  
7f275c3cb000-7f275c3cc000 rw-p 00009000 103:05 923315  
/usr/lib/x86\_64-linux-gnu/libacl.so.1.1.2253  
7f275c3cc000-7f275c3cf000 r--p 00000000 103:05 923446  
/usr/lib/x86\_64-linux-gnu/libcanberra.so.0.2.5  
7f275c3cf000-7f275c3d9000 r-xp 00003000 103:05 923446  
/usr/lib/x86\_64-linux-gnu/libcanberra.so.0.2.5  
7f275c3d9000-7f275c3dd000 r--p 0000d000 103:05 923446  
/usr/lib/x86\_64-linux-gnu/libcanberra.so.0.2.5  
7f275c3dd000-7f275c3de000 r--p 00010000 103:05 923446  
/usr/lib/x86\_64-linux-gnu/libcanberra.so.0.2.5  
7f275c3de000-7f275c3df000 rw-p 00011000 103:05 923446  
/usr/lib/x86\_64-linux-gnu/libcanberra.so.0.2.5  
7f275c3df000-7f275c3e5000 r--p 00000000 103:05 924431  
/usr/lib/x86\_64-linux-gnu/libselinux.so.1  
7f275c3e5000-7f275c3fe000 r-xp 00006000 103:05 924431  
/usr/lib/x86\_64-linux-gnu/libselinux.so.1  
7f275c3fe000-7f275c405000 r--p 0001f000 103:05 924431

```

/usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c405000-7f275c406000 ---p 00026000 103:05 924431
/usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c406000-7f275c407000 r--p 00026000 103:05 924431
/usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c407000-7f275c408000 rw-p 00027000 103:05 924431
/usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c408000-7f275c40a000 rw-p 00000000 00:00 0
7f275c40a000-7f275c418000 r--p 00000000 103:05 924540
/usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c418000-7f275c427000 r-xp 0000e000 103:05 924540
/usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c427000-7f275c435000 r--p 0001d000 103:05 924540
/usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c435000-7f275c439000 r--p 0002a000 103:05 924540
/usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c439000-7f275c43a000 rw-p 0002e000 103:05 924540
/usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c43a000-7f275c449000 r--p 00000000 103:05 917895
/usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c449000-7f275c4f0000 r-xp 0000f000 103:05 917895
/usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c4f0000-7f275c587000 r--p 000b6000 103:05 917895
/usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c587000-7f275c588000 r--p 0014c000 103:05 917895
/usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c588000-7f275c589000 rw-p 0014d000 103:05 917895
/usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c589000-7f275c58b000 rw-p 00000000 00:00 0
7f275c5ae000-7f275c5af000 r--p 00000000 103:05 917889
/usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5af000-7f275c5d2000 r-xp 00001000 103:05 917889
/usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5d2000-7f275c5da000 r--p 00024000 103:05 917889
/usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5db000-7f275c5dc000 r--p 0002c000 103:05 917889
/usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5dc000-7f275c5dd000 rw-p 0002d000 103:05 917889
/usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5dd000-7f275c5de000 rw-p 00000000 00:00 0
7ffd22d2f000-7ffd22d50000 rw-p 00000000 00:00 0 [stack]
7ffd22db0000-7ffd22db4000 r--p 00000000 00:00 0 [vvar]
7ffd22db4000-7ffd22db6000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0
[vsyscall]

```

True	False
------	-------



This is a virtual memory map (not physical memory map)



The size of the stack is one page

True	False	
<input type="radio"/>	<input type="radio"/>	vim.basic uses the math library
<input type="radio"/>	<input type="radio"/>	The size of the heap is one page
<input type="radio"/>	<input type="radio"/>	The 5th entry 55a4352c5000-55a4352e2000 <b>may</b> correspond to "data" of the vim.basic

This is a virtual memory map (not physical memory map): True

The size of the stack is one page: False

vim.basic uses the math library: True

The size of the heap is one page: False

The 5th entry 55a4352c5000-55a4352e2000 **may** correspond to "data" of the vim.basic: True

#### Question 15

Not answered

Marked out of 0.50

Map the parts of a C code to the memory regions they are related to

local variables	<input data-bbox="649 850 797 903" type="text" value="Choose..."/>
function arguments	<input data-bbox="649 913 797 966" type="text" value="Choose..."/>
global un-initialized variables	<input data-bbox="649 976 797 1029" type="text" value="Choose..."/>
malloced memory	<input data-bbox="649 1039 797 1092" type="text" value="Choose..."/>
functions	<input data-bbox="649 1102 797 1155" type="text" value="Choose..."/>
global initialized variables	<input data-bbox="649 1165 797 1218" type="text" value="Choose..."/>
static variables	<input data-bbox="649 1228 797 1281" type="text" value="Choose..."/>

The correct answer is: local variables → stack, function arguments → stack, global un-initialized variables → bss, malloced memory → heap, functions → code, global initialized variables → data, static variables → data

#### Question 16

Not answered

Marked out of 1.00

Shared memory is possible with which of the following memory management schemes ?

Select one or more:

- ☐ a. continuous memory management
- ☐ b. segmentation
- ☐ c. paging
- ☐ d. demand paging

The correct answers are: paging, segmentation, demand paging



[◀ \(Code\) mmap related programs](#)

[Jump to...](#)

[Points from Mid-term feedback ▶](#)